

Movielens
Movie Recommendation System
A Harvard Capstone Project

Manoj Bijoor

July 28, 2021

Abstract

Develop an algorithm using the edx set (MovieLens data). For a final test of my final algorithm, predict movie ratings in the validation set (the final hold-out test set) as if they were unknown. RMSE is to be used to evaluate how close my predictions are to the true values in the validation set (the final hold-out test set). My target is $\text{RMSE} < 0.86490$.

I achieved my goal with an RMSE of 0.8634961.

Contents

List of tables	v
List of figures	vi
List of Equations	vii
1 Project Overview: MovieLens - A Harvard Capstone Project	1
1.1 Create Train and Final Hold-out Test Sets	1
1.1.1 Important: Data sets usage	1
1.2 Final Product	1
1.2.1 My submission for this project is three files:	1
2 Exploratory Data Analysis	3
2.1 Data Wrangling	3
2.1.1 Initial data Exploration & Visualization	3
2.1.1.1 A Very Sparse Matrix	5
2.1.1.2 General Properties of the data	6
2.1.2 Further data Exploration, Visualization & Modification	8
2.1.2.1 Modify edx	8
2.1.2.2 Modify validation, repeat above steps	9
2.1.2.3 Genres combinations per movie - a closer look	10
2.1.2.4 Movie Rating Date-Time - a closer look	11
2.1.2.5 Movie Release Date - a closer look	13
3 Analysis - Model Building and Evaluation	20
3.1 Split the edx data into separate training and test sets	20
3.1.1 Loss function	20
3.2 Model 1: A first naive “mean” model	21
3.2.1 Results Model 1	22
3.3 Model 2: Movie effects	23
3.3.1 Results Model 1-2	25
3.4 Model 3: User effects	26
3.4.1 Results Table Model 1-3	30
3.5 Model 4: Genre effects	31
3.5.1 Results Table Model 1-4	34
3.6 Model 5: Rating Time effect	35
3.6.1 Results Table Model 1-5	37
3.7 Model 6: Release Date Effect	38
3.7.1 Results Table Model 1-6	40
3.8 Regularization	41
3.8.1 Motivation	41
3.8.2 Penalized least squares	44
3.9 Model 7 : Model 2 + Regularization: Choosing the penalty terms	45
3.9.1 Effect of use of penalized estimates	47
3.9.2 Results Table Model 1-7	49
3.10 Model 8 : Model 4 + Regularization: Choosing the penalty terms	50
3.10.1 Results Table Model 1-8	52
3.11 Model 9 : Model 6 + Regularization: Choosing the penalty terms	53
3.11.1 Results Table Model 1-9	56
3.12 Multi-fold cross-validation on only edx train_set data set, for the full model 9.	57
3.13 Model 10 : Model 6 + Regularization from multi-fold CV(Model 9 with multi-fold CV)	62
3.13.1 Results Table Model 1-10	64

3.14 Final Hold-Out test of my final algorithm	65
3.14.1 Results Table Model 1-10	67
4 Conclusion	68
5 Appendix: All code for this report	69
Alphabetical Index	91

List of tables

1	Movielens data	3
2	Unique Users, Movies and Genres	3
3	Matrix of seven users and four movies	4
4	Movielens edx data with rating date-time	8
5	Movielens edx data with movie release date	8
6	Movielens validation data with rating date-time	9
7	Movielens validation data with movie release date	9
8	Movielens edx data with average rating due to rating time effect	11
9	Movielens validation data with average rating due to rating time effect	12
10	Movielens edx data with average rating due to release date effect	17
11	Movielens validation data with average rating due to release date effect	19
12	RMSE Results Model 1	22
13	RMSE Results Models 1-2	25
14	RMSE Results Models 1-3	30
15	RMSE Results Models 1-4	34
16	RMSE Results Models 1-5	37
17	RMSE Results Models 1-6	40
18	Without Regularization 10 largest mistakes	41
19	Without Regularization 10 Best movies	42
20	Without Regularization 10 Worst movies	42
21	With Regularization top 10 best movies based on the penalized estimates	47
22	With Regularization top 10 worst movies based on the penalized estimates	48
23	RMSE Results Models 1-7	49
24	RMSE Results Models 1-8	52
25	RMSE Results Models 1-9	56
26	RMSE Results Models 1-10	64
27	RMSE Results Models 1-10 Validation	67

List of figures

1	A Very Sparse Matrix	5
2	Movies getting rated distribution	6
3	Users rating movies distribution	7
4	Movies genres error bar plots	10
5	Movies average ratings for each week versus day	11
6	Ratings Movie Release Date - All dates	13
7	25 Movies with the most ratings per year and their average rating post 1993	14
8	Movies average ratings versus ratings per year post 1993	15
9	Movies average ratings versus ratings per year pre 1993	16
10	Movies average ratings versus ratings per year for all years for edx	17
11	Movie effect or bias distribution	24
12	Average rating for users who have rated over 100 movies	26
13	Average rating for users who have rated any movies	27
14	User effect or bias distribution	29
15	Average rating for movies of category genres	31
16	Genres effect or bias distribution	33
17	Rating time effect or bias distribution	36
18	Release Date effect or bias distribution	39

List of Equations

1	Movie Rating Date-Time effect Equation 1	11
2	Loss function RMSE Equation 2	20
3	Model 1: A naive "mean" model Equation 3	21
4	Model 2: Movie effects linear model Equation 4	23
5	LSE linear function to fit Movie effects linear model Equation 5	23
6	Movie specific effects Equation 6	23
7	Model 3: Movie + User effects linear model Equation 7	27
8	LSE linear function to fit Movie + User effects linear model Equation 8	27
9	User specific effects Equation 9	27
10	Model 4: Movie + User + Genre effects linear model Equation 10	31
11	LSE linear function to fit Movie + User + Genres effects linear model Equation 11	31
12	Genres specific effects Equation 12	32
13	Genres specific effects Equation 13	32
14	Model 5: Movie + User + Genre + Rating time effects linear model Equation 14	35
15	LSE linear function to fit Movie + User + Genres + Rating time effects linear model Equation 15	35
16	Rating time specific effects Equation 16	35
17	Rating time specific effects in function form Equation 17	35
18	Model 6: Movie + User + Genre + Rating time + Release date effects linear model Equation 18	38
19	LSE linear function to fit Movie + User + Genres + Rating time + Release date effects linear model Equation 19	38
20	Release date specific effects Equation 20	38
21	Release date specific effects in function form Equation 21	38

1 Project Overview: MovieLens - A Harvard Capstone Project

A movie recommendation system using the MovieLens dataset.

For this project, I will be creating a movie recommendation system using the MovieLens dataset, provided by GroupLens Research¹, a research lab in the Department of Computer Science and Engineering at the University of Minnesota, Twin Cities specializing in recommender systems, online communities, mobile and ubiquitous technologies, digital libraries, and local geographic information systems.

GroupLens Research² has collected and made available rating data sets from the MovieLens web site³. The data sets were collected over various periods of time, depending on the size of the set.

I will use the 10M version of the MovieLens dataset⁴ to make the computation a little easier.

First, I will download the MovieLens data and run code provided to generate my datasets.

Second, I will train a machine learning algorithm using the inputs in one subset to predict movie ratings in the validation set.

1.1 Create Train and Final Hold-out Test Sets

I will develop my algorithm using the edx set. For a final test of my final algorithm, I predict movie ratings in the validation set (the final hold-out test set) as if they were unknown. RMSE⁵ will be used to evaluate how close my predictions are to the true values in the validation set (the final hold-out test set). My target is RMSE < 0.86490.

1.1.1 Important: Data sets usage

The validation data (the final hold-out test set) will NOT be used for training, developing, or selecting my algorithm and it will ONLY be used for evaluating the RMSE of my final algorithm. The final hold-out test set will only be used at the end of my project with my final model. It will not be used to test the RMSE of multiple models during model development. I will split the edx data into separate training and test sets to design and test my algorithm.

1.2 Final Product

1.2.1 My submission for this project is three files:

1. My report in Rmd format
2. My report in PDF format (knit from my Rmd file)
3. A script in R format that generates my predicted movie ratings and RMSE score (contains all code and comments for my project)

The report documents the analysis and presents the findings, along with supporting statistics and figures. The report assumes that the reader is not familiar with the project or the data. The report includes the RMSE generated and the following sections:

1. an introduction/overview/executive summary section that describes the dataset and summarizes the goal of the project and key steps that were performed

¹<https://grouplens.org/>

²<https://grouplens.org/datasets/movielens/>

³<https://movielens.org>

⁴<https://grouplens.org/datasets/movielens/10m/>

⁵https://en.wikipedia.org/wiki/Root-mean-square_deviation

2. a methods/analysis section that explains the process and techniques used, including data cleaning, data exploration and visualization, insights gained, and my modeling approach
3. a results section that presents the modeling results and discusses the model performance
4. a conclusion section that gives a brief summary of the report, its limitations and future work

2 Exploratory Data Analysis

2.1 Data Wrangling

Data wrangling⁶, sometimes referred to as data munging, is the process of transforming and mapping data from one “raw” data form into another format with the intent of making it more appropriate and valuable for a variety of downstream purposes such as analytics.

The main steps could be described as follows:

1. Discovering
2. Structuring
3. Cleaning
4. Enriching
5. Validating
6. Publishing

Let’s perform the steps or combinations thereof starting with Initial data Exploration & Visualization in the next few subsections.

2.1.1 Initial data Exploration & Visualization

The first ten rows out of 9000055 rows of the Movielens data can be found in Table 1.

Table 1: Movielens data

userId	movieId	rating	timestamp	title	genres
1	122	5	838985046	Boomerang (1992)	Comedy Romance
1	185	5	838983525	Net, The (1995)	Action Crime Thriller
1	292	5	838983421	Outbreak (1995)	Action Drama Sci-Fi Thriller
1	316	5	838983392	Stargate (1994)	Action Adventure Sci-Fi
1	329	5	838983392	Star Trek: Generations (1994)	Action Adventure Drama Sci-Fi
1	355	5	838984474	Flintstones, The (1994)	Children Comedy Fantasy
1	356	5	838983653	Forrest Gump (1994)	Comedy Drama Romance War
1	362	5	838984885	Jungle Book, The (1994)	Adventure Children Romance
1	364	5	838983707	Lion King, The (1994)	Adventure Animation Children Drama Musical
1	370	5	838984596	Naked Gun 33 1/3: The Final Insult (1994)	Action Comedy

Each row represents a rating given by one user to one movie.

We can see the number of unique users that provided ratings and how many unique movies were rated. The unique genres here is based on a column called “genres” that includes every genre that applies to the movie. Some movies fall under several genres. Defining a category as whatever combination appears in this column, we are going to refer to this category as “unique genres” or simply “genres” from now on.

The number of unique users, movies and genres can be found in Table 2.

Table 2: Unique Users, Movies and Genres

unique_users	unique_movies	unique_genres
69878	10677	797

If we multiply the number of unique users by number of unique movies, we get a very large number, actually 746087406, yet our data table has 9000055 rows. This implies that not every user rated every movie. So

⁶https://en.wikipedia.org/wiki/Data_wrangling

we can think of these data as a very large matrix, with users on the rows and movies on the columns, with many empty cells. The ‘gather’ function permits us to convert it to this format, but if we try it for the entire matrix, it will crash R.

Let’s show the matrix of seven users ie: userId’s 13-20 and four movies in Table 3.

Table 3: Matrix of seven users and four movies

userId	Forrest Gump (1994)	Jurassic Park (1993)	Pulp Fiction (1994)	Silence of the Lambs, The (1991)
13	NA	NA	4	NA
16	NA	3	NA	NA
17	NA	NA	NA	5
18	NA	3	5	5
19	4	1	NA	NA

2.1.1.1 A Very Sparse Matrix You can think of the task of a recommendation system as filling in the ‘NAs’ in the table above. To see how sparse the matrix is, here is the matrix in Figure 1 for a random sample of 100 movies and 100 users with yellow indicating a user/movie combination for which we have a rating.

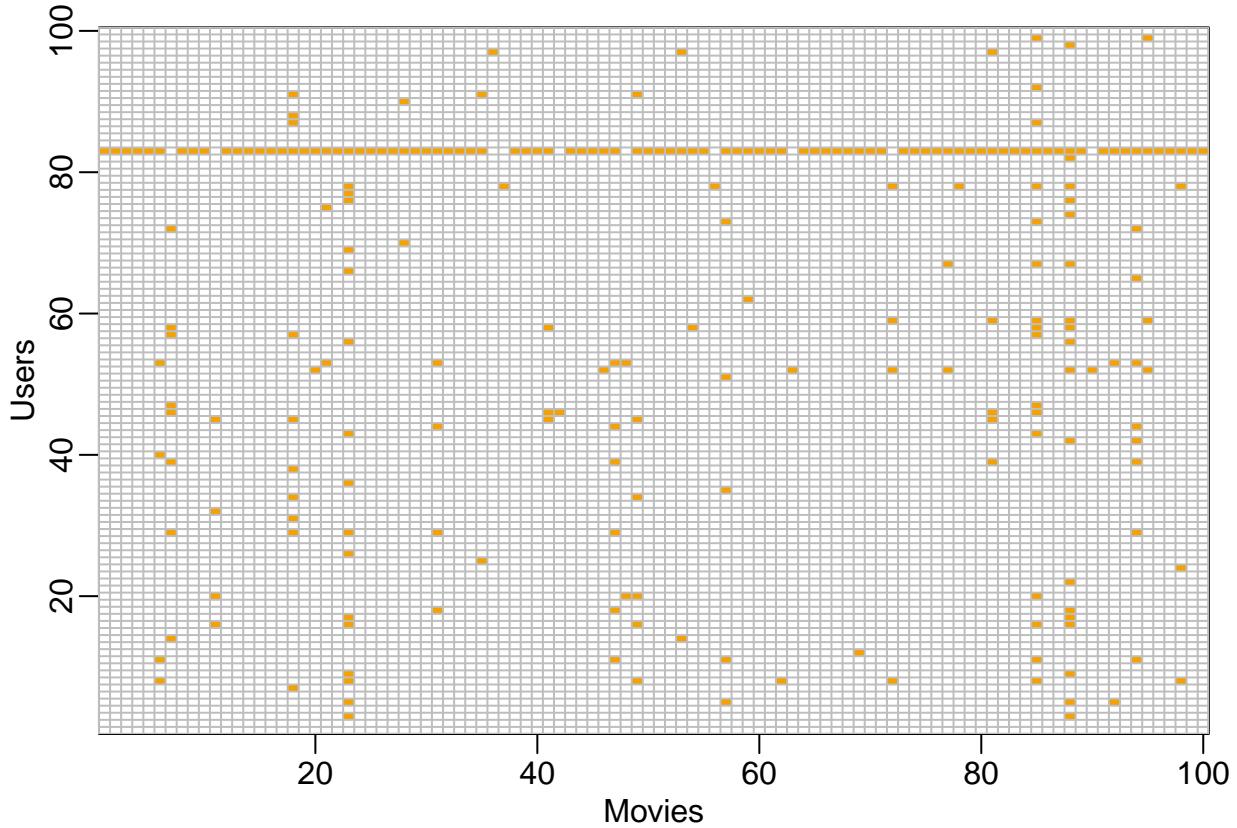


Figure 1: A Very Sparse Matrix

This machine learning challenge is quite complicated, because each outcome Y has a different set of predictors. To see this, note that if we are predicting the rating for movie i by user u , in principle, all other ratings related to movie i and by user u may be used as predictors, but different users rate different movies and a different number of movies. Furthermore, we may be able to use information from other movies that we have determined are similar to movie i or from users determined to be similar to user u . In essence, the entire matrix can be used as predictors for each cell.

2.1.1.2 General Properties of the data Let's look at some of the *general properties* of the data to better understand the challenges.

The *first thing* we notice is that some movies get rated more than others. Figure 2 shows the Movies getting rated distribution. This should not surprise us given that there are blockbuster movies watched by millions and artsy, independent movies watched by just a few:

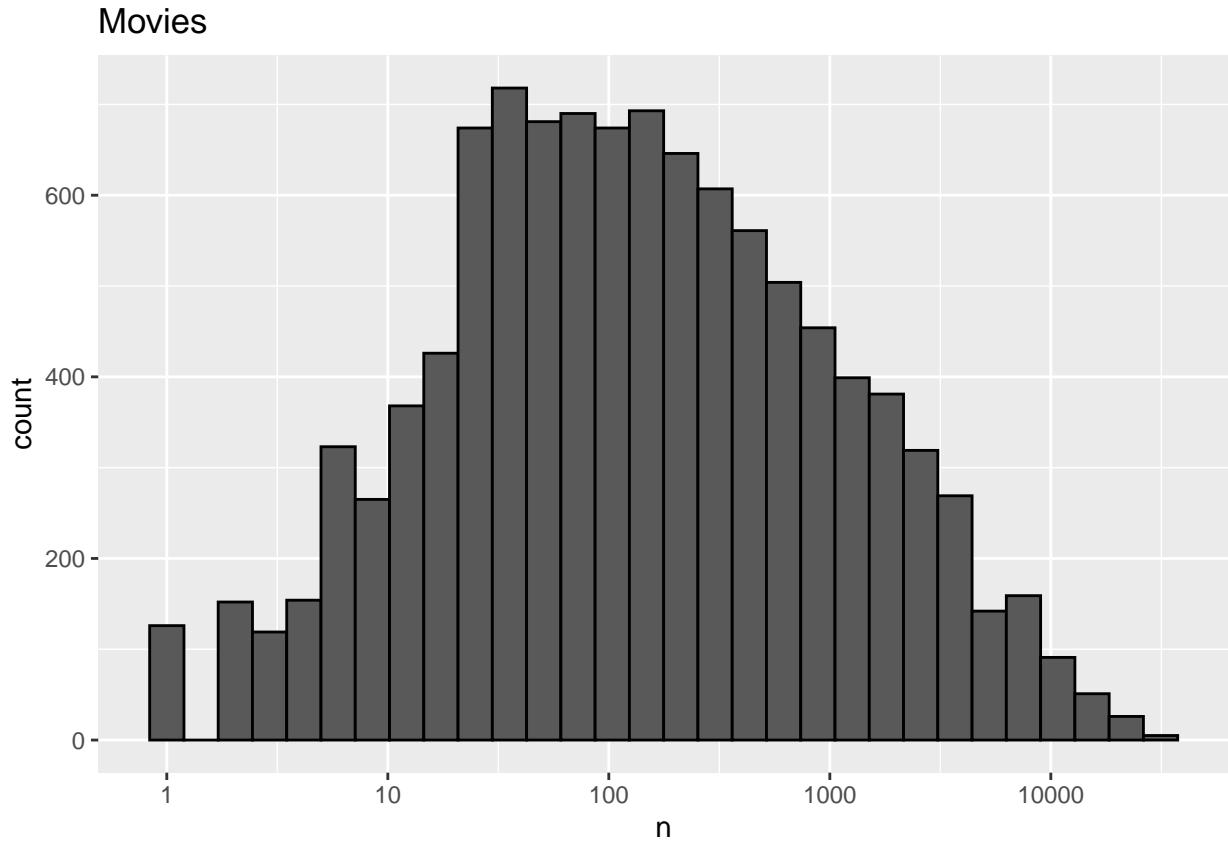


Figure 2: Movies getting rated distribution

Our ***second observation*** is that some users are more active than others at rating movies. Figure 3 shows Users rating movies distribution:

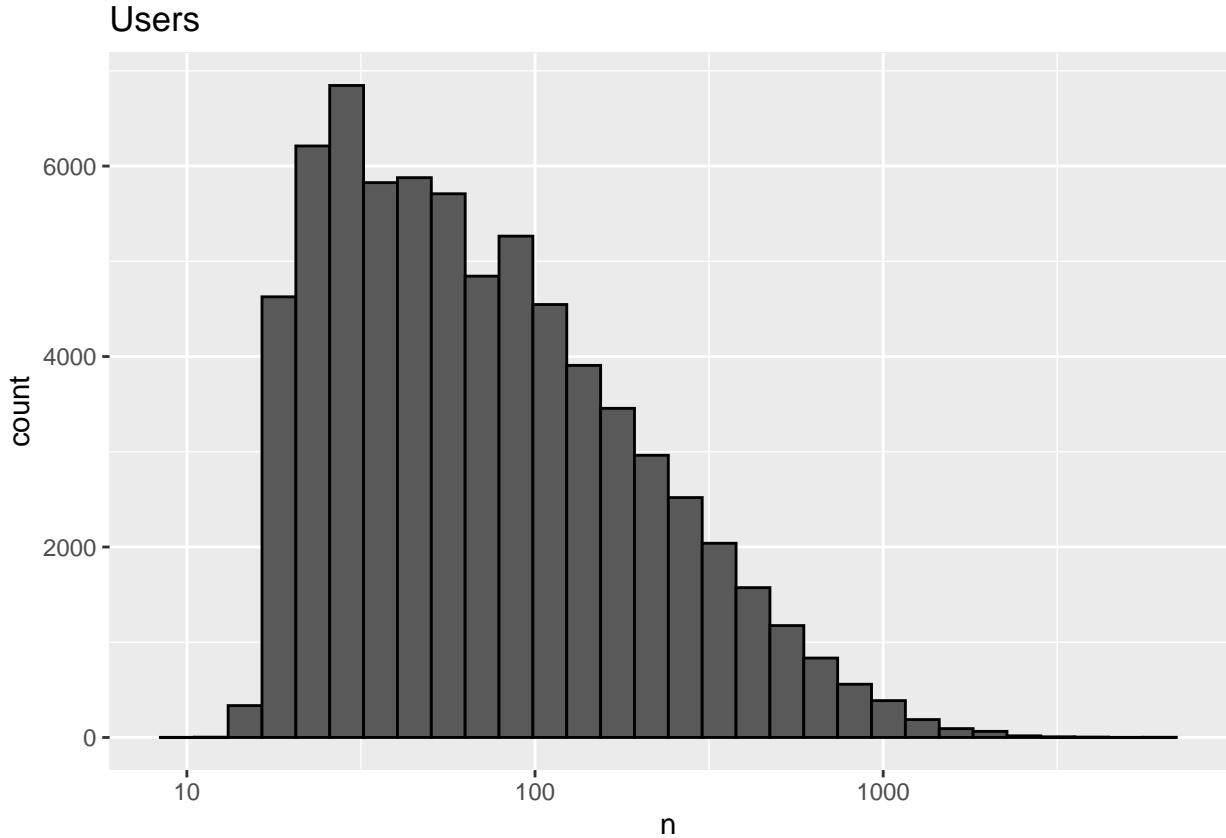


Figure 3: Users rating movies distribution

2.1.2 Further data Exploration, Visualization & Modification

2.1.2.1 Modify edx Convert timestamp in Movielens edx data Table 1 into date-time, a more readable and useful format named *rating_date* in Table 4 below.

Table 4: Movielens edx data with rating date-time

userId	movieId	rating	title	genres	rating_date
1	122	5	Boomerang (1992)	Comedy Romance	1996-08-02 11:24:06
1	185	5	Net, The (1995)	Action Crime Thriller	1996-08-02 10:58:45
1	292	5	Outbreak (1995)	Action Drama Sci-Fi Thriller	1996-08-02 10:57:01
1	316	5	Stargate (1994)	Action Adventure Sci-Fi	1996-08-02 10:56:32
1	329	5	Star Trek: Generations (1994)	Action Adventure Drama Sci-Fi	1996-08-02 10:56:32

Split title in Movielens edx data Table 1 into title and year movie released, a more useful format named *movie_dt* in Table 5 below.

Table 5: Movielens edx data with movie release date

userId	movieId	rating	title	genres	rating_date	movie_dt
1	122	5	Boomerang	Comedy Romance	1996-08-02 11:24:06	1992
1	185	5	Net, The	Action Crime Thriller	1996-08-02 10:58:45	1995
1	292	5	Outbreak	Action Drama Sci-Fi Thriller	1996-08-02 10:57:01	1995
1	316	5	Stargate	Action Adventure Sci-Fi	1996-08-02 10:56:32	1994
1	329	5	Star Trek: Generations	Action Adventure Drama Sci-Fi	1996-08-02 10:56:32	1994

2.1.2.2 Modify validation, repeat above steps Convert timestamp in MovieLens validation data Table 1 into date-time, a more readable and useful format named *rating_date* in Table 6 below.

Table 6: MovieLens validation data with rating date-time

userId	movieId	rating	title	genres	rating_date
1	231	5	Dumb & Dumber (1994)	Comedy	1996-08-02 10:56:32
1	480	5	Jurassic Park (1993)	Action Adventure Sci-Fi Thriller	1996-08-02 11:00:53
1	586	5	Home Alone (1990)	Children Comedy	1996-08-02 11:07:48
2	151	3	Rob Roy (1995)	Action Drama Romance War	1997-07-07 03:34:10
2	858	2	Godfather, The (1972)	Crime Drama	1997-07-07 03:20:45

Split title in MovieLens validation data Table 1 into title and year movie released, a more useful format named *movie_dt* in Table 7 below.

Table 7: MovieLens validation data with movie release date

userId	movieId	rating	title	genres	rating_date	movie_dt
1	231	5	Dumb & Dumber	Comedy	1996-08-02 10:56:32	1994
1	480	5	Jurassic Park	Action Adventure Sci-Fi Thriller	1996-08-02 11:00:53	1993
1	586	5	Home Alone	Children Comedy	1996-08-02 11:07:48	1990
2	151	3	Rob Roy	Action Drama Romance War	1997-07-07 03:34:10	1995
2	858	2	Godfather, The	Crime Drama	1997-07-07 03:20:45	1972

2.1.2.3 Genres combinations per movie - a closer look The movielens data Table 8 also has a genres column. This column includes every genre that applies to the movie. Some movies fall under several genres. We define a category of genres as whatever combination of genres appears in this column, and refer to it as simply “genres”.

Here we keep only categories with more than 1,000 ratings. Then compute the average and standard error for each category, and plot these as error bar plots. See Figure 4 :

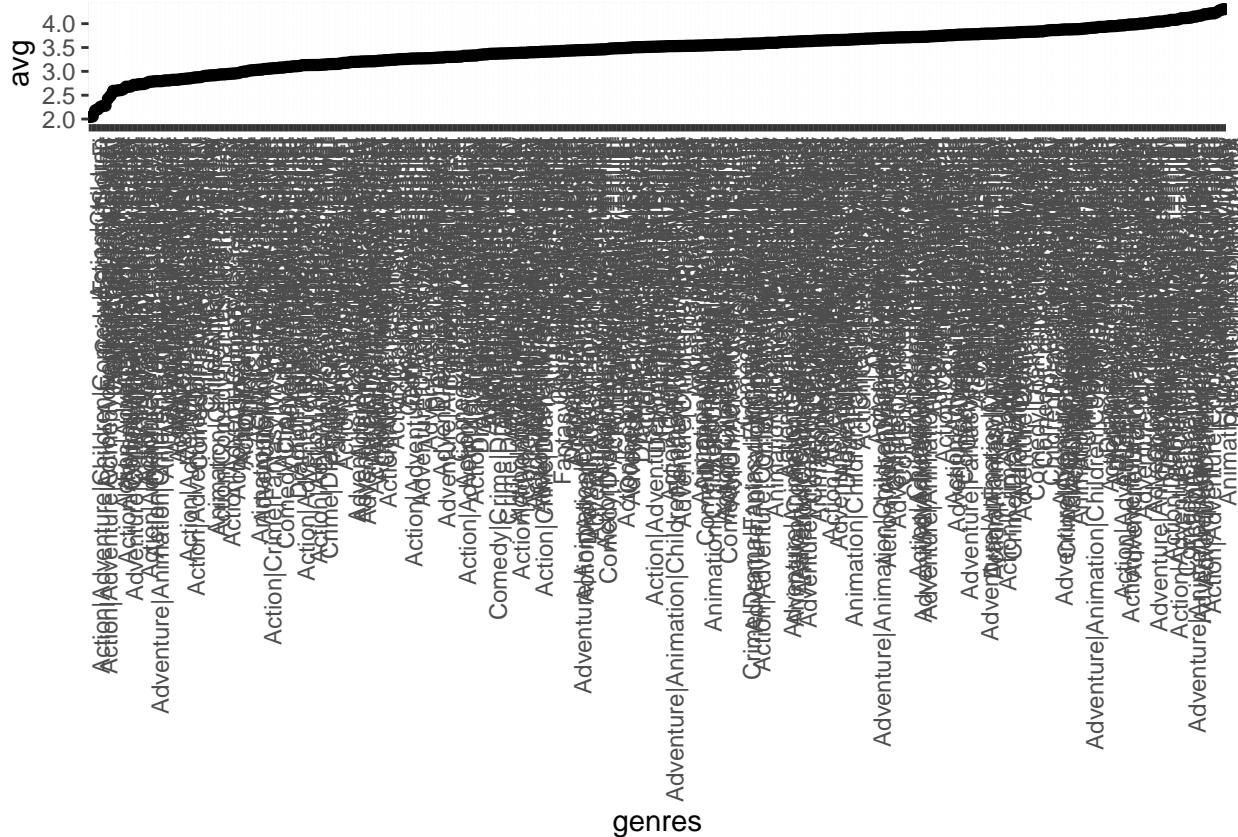


Figure 4: Movies genres error bar plots

The plot shows strong evidence of a genre effect.

2.1.2.4 Movie Rating Date-Time - a closer look The Movielens edx data Table 1 also includes a time stamp. This variable represents the time and date in which the rating was provided. The units are seconds since January 1, 1970. We create a new column date with the date named *rating_date* in subsection [Modify edx](#) to get Table 5 .

We compute the average rating for each week and plot this average against day. See Figure 5 :

```
'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

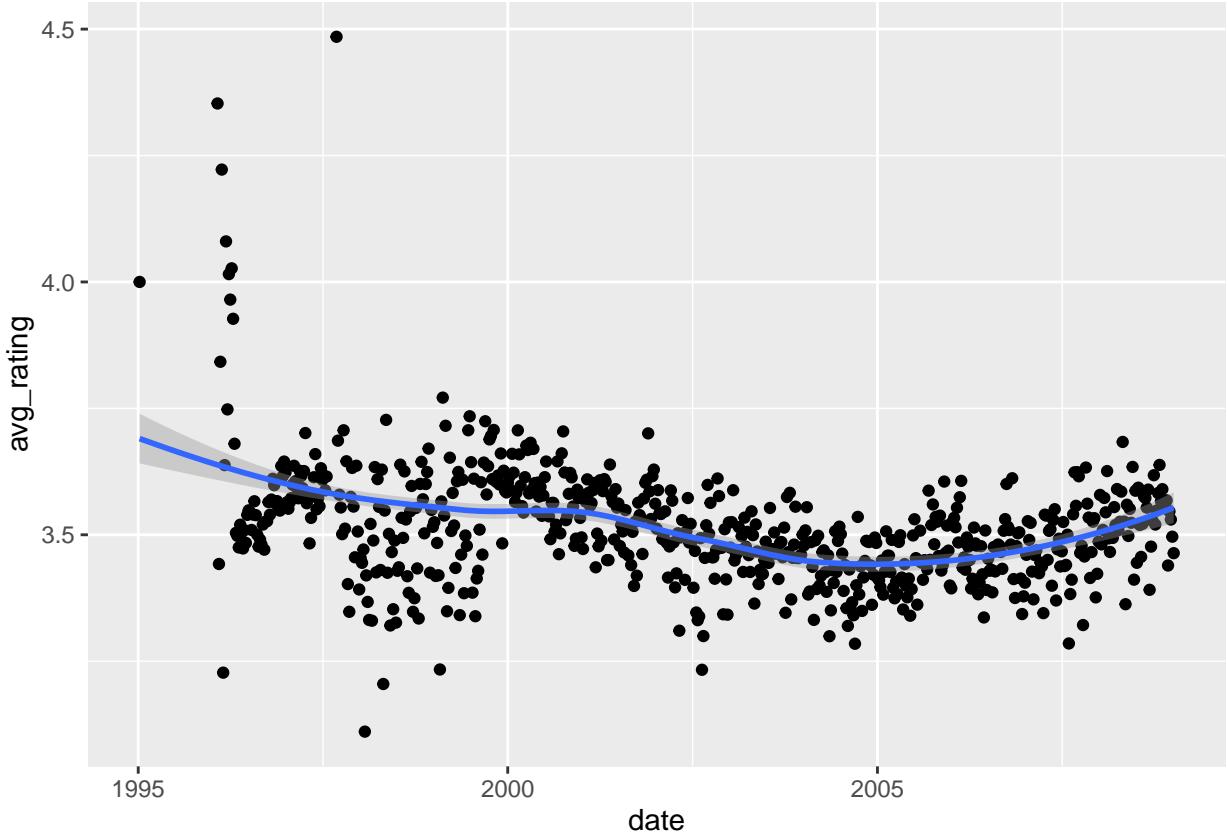


Figure 5: Movies average ratings for each week versus day

The plot shows some evidence of a time effect. If we define $d_{u,i}$ as the day for user's u rating of movie i , then the following model given by Equation 1 is most appropriate:

$$Y_{u,i} = \mu + b_i + b_u + f(d_{u,i}) + \epsilon_{u,i}, \text{ with } f \text{ a smooth function of } d_{u,i} \quad (1)$$

Modify edx Let's update the *edx* table with a new column for the average rating for each week and another column for the day rounded to the nearest value of the week to get Table 8 below:

Table 8: Movielens edx data with average rating due to rating time effect

userId	movieId	rating	title	genres	rating_date	movie_dt	date	avg_rating
1	122	5	Boomerang	Comedy Romance	1996-08-02 11:24:06	1992	1996-08-04	3.538801
1	185	5	Net, The	Action Crime Thriller	1996-08-02 10:58:45	1995	1996-08-04	3.538801
1	292	5	Outbreak	Action Drama Sci-Fi Thriller	1996-08-02 10:57:01	1995	1996-08-04	3.538801
1	316	5	Stargate	Action Adventure Sci-Fi	1996-08-02 10:56:32	1994	1996-08-04	3.538801
1	329	5	Star Trek: Generations	Action Adventure Drama Sci-Fi	1996-08-02 10:56:32	1994	1996-08-04	3.538801

TODO: Repeat above for validation data as well and somehow add this to the modelling section

Modify validation We need to do the above *avg_rating_time_effect* update for the validation data as well. Let's update the *validation* table to get Table 9 below:

Table 9: MovieLens validation data with average rating due to rating time effect

userId	movieId	rating	title	genres	rating_date	movie_dt	date	avg_rating
1	231	5	Dumb & Dumber	Comedy	1996-08-02 10:56:32	1994	1996-08-04	3.555820
1	480	5	Jurassic Park	Action Adventure Sci-Fi Thriller	1996-08-02 11:00:53	1993	1996-08-04	3.555820
1	586	5	Home Alone	Children Comedy	1996-08-02 11:07:48	1990	1996-08-04	3.555820
2	151	3	Rob Roy	Action Drama Romance War	1997-07-07 03:34:10	1995	1997-07-06	3.606571
2	858	2	Godfather, The	Crime Drama	1997-07-07 03:20:45	1972	1997-07-06	3.606571

2.1.2.5 Movie Release Date - a closer look Computing the number of ratings for each movie and then plotting it against the year the movie came out, that is the release date and using the square root transformation on the counts using Table 5 , we get see Figure 6 :

TODO: Align images

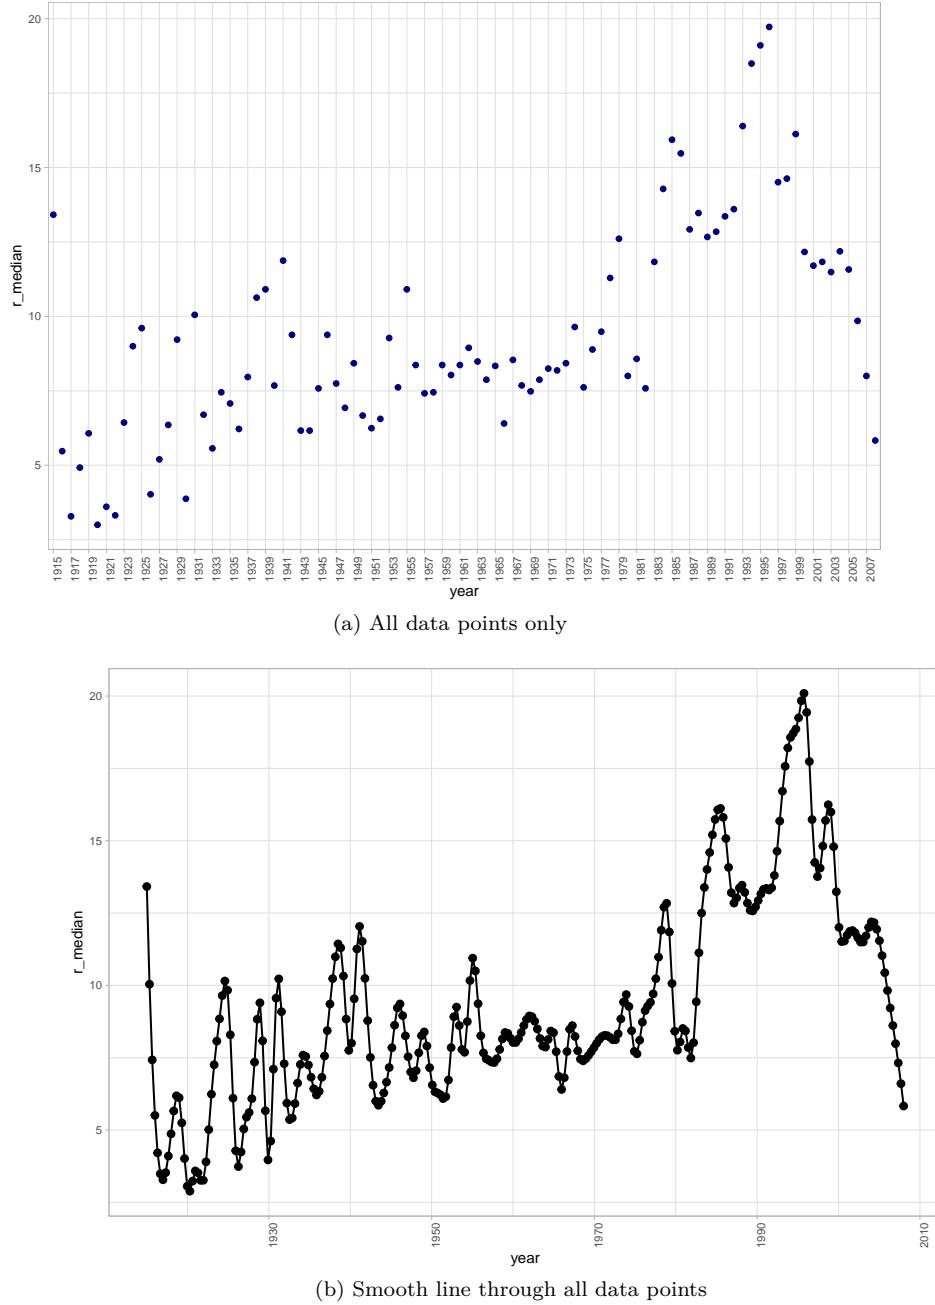


Figure 6: Ratings Movie Release Date - All dates

we see that, on average, movies that came out after 1993 get more ratings. We also see that with newer movies, starting in 1993, the number of ratings decreases with year: the more recent a movie is, the less time users have had to rate it.

Among movies that came out in 1993 or later, we select the top 25 movies with the highest average number of ratings per year (n/year) and calculate the average rating of each of them. To calculate number of ratings per year, use 2018 as the end year. See Figure 7 :

```
'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```

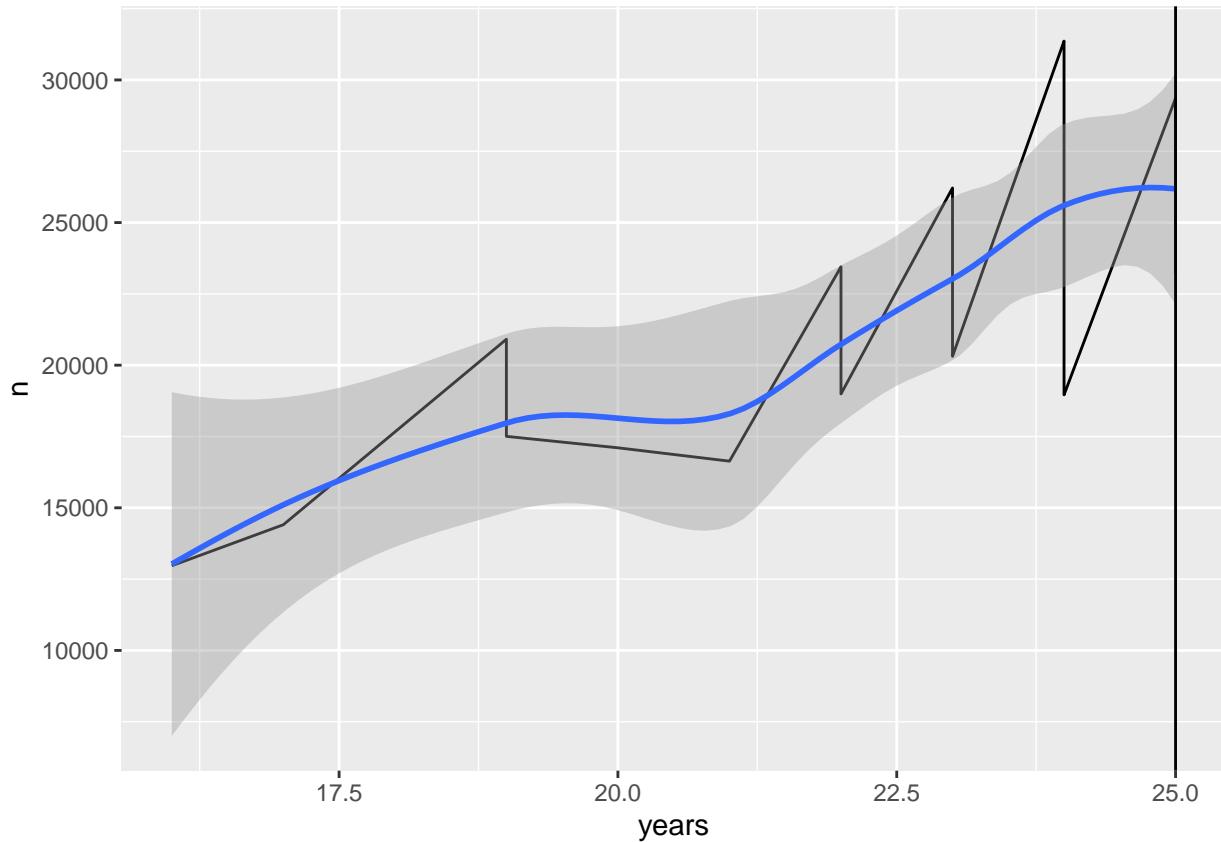


Figure 7: 25 Movies with the most ratings per year and their average rating post 1993

We see that the most rated movies tend to have above average ratings. This is not surprising: more people watch popular movies. To confirm this, we stratify the post 1993 movies by ratings per year and compute their average ratings. Figure 8 is a plot of average ratings versus ratings per year showing an estimate of the trend.

We see that the more a movie is rated, the higher the rating.

Post-1993 movies

```
'geom_smooth()' using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

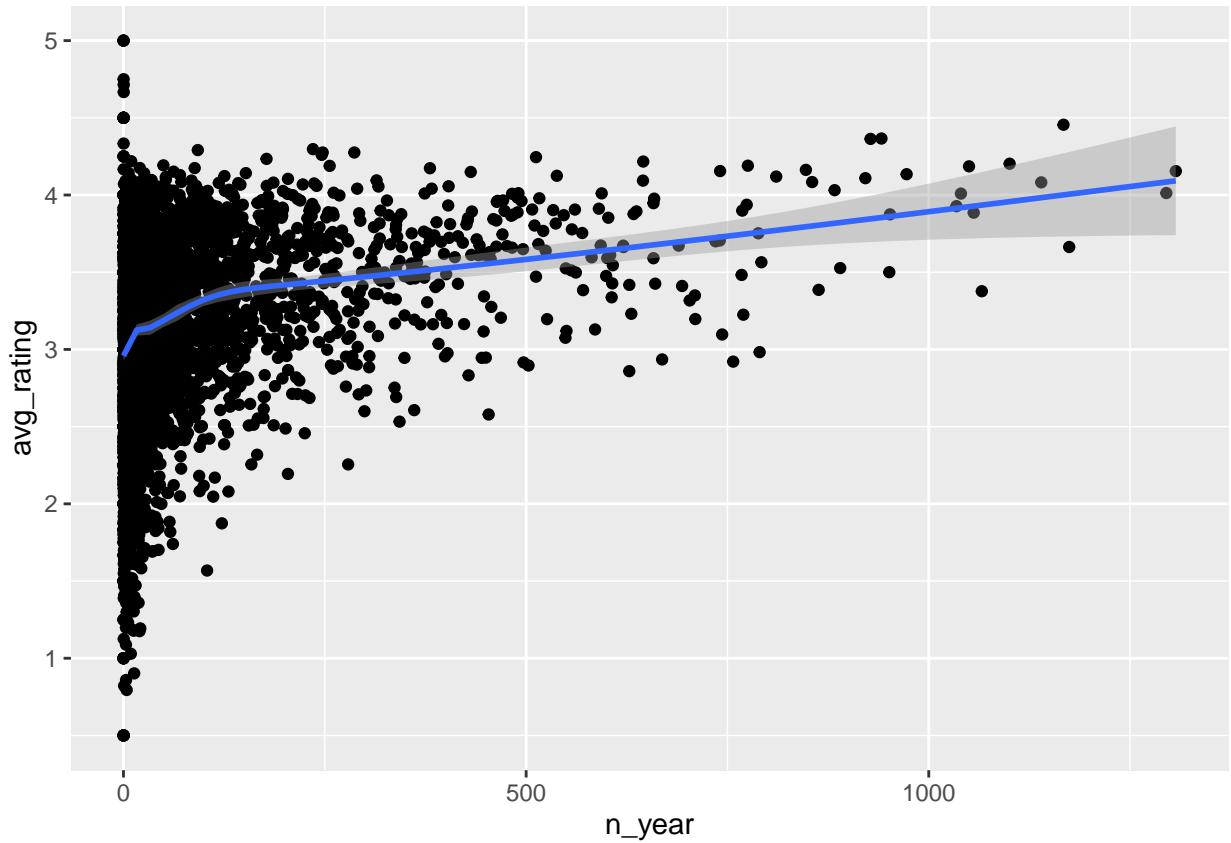


Figure 8: Movies average ratings versus ratings per year post 1993

Pre-1993 movies

Compare Pre-1993 movies trend shown here in Figure 9 Versus Post-1993 movies trend in Figure 8 above.

```
'geom_smooth()' using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

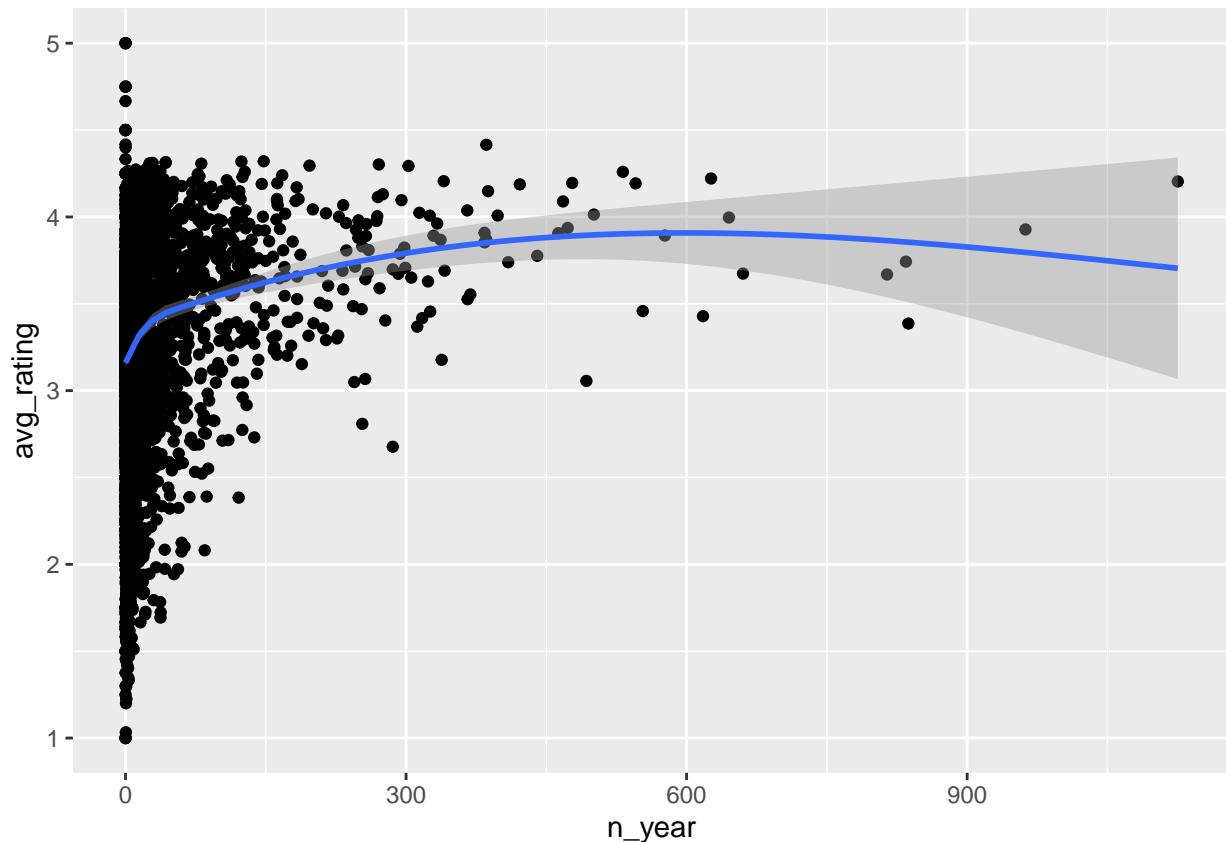


Figure 9: Movies average ratings versus ratings per year pre 1993

Modify edx data for Release Date Effect We stratify the movies by ratings per year and compute their average ratings based on what we learnt above, where we confirmed our intuition that more people watch popular movies. Finally edx data table looks as shown in Table 10 below. Figure 10 is a plot of average ratings versus ratings per year showing an estimate of the trend.

All Years

```
'geom_smooth()' using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

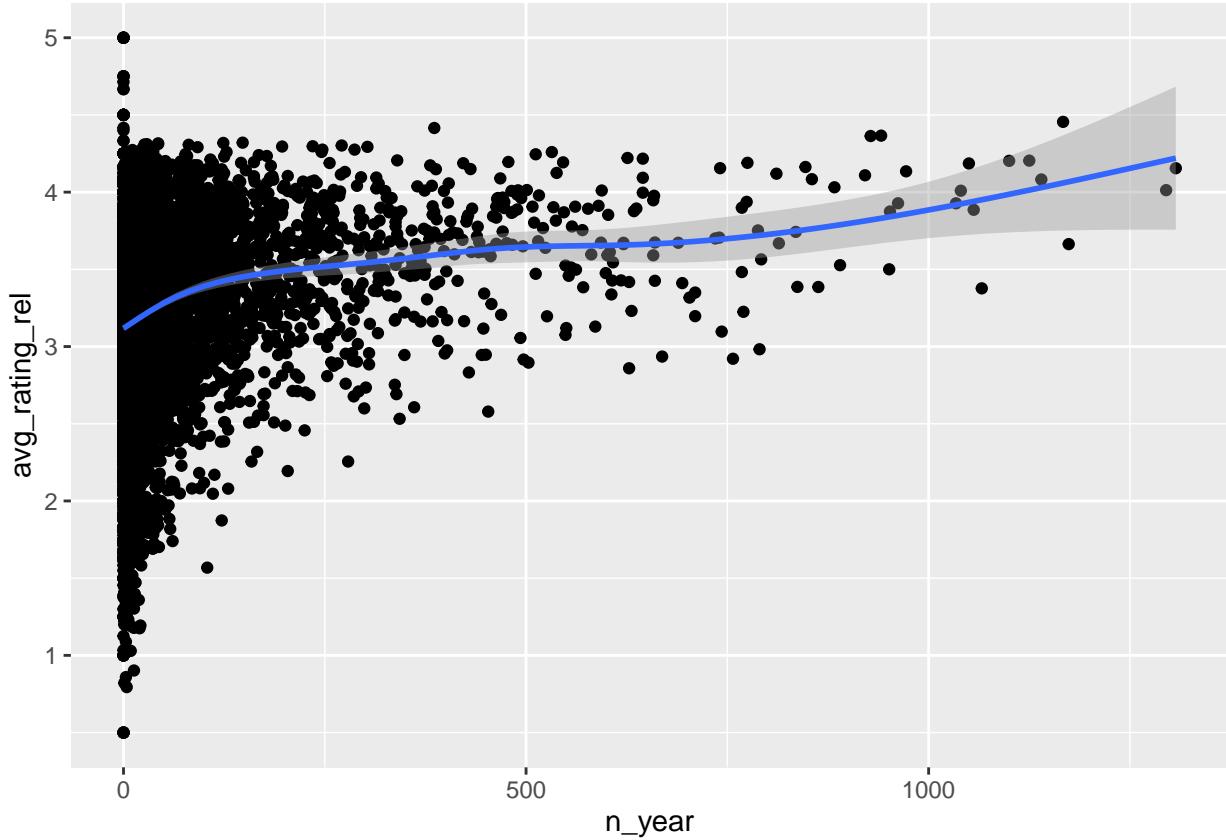


Figure 10: Movies average ratings versus ratings per year for all years for edx

Table 10: MovieLens edx data with average rating due to release date effect

userId	movieId	rating	title	genres	rating_date	movie_dt	date	avg_rating	avg_rating_rel	n	years	n_year
1	122	5	Boomerang	Comedy Romance	1996-08-02 11:24:06	1992	1996-08-04	3.538801	2.858586	2178	26	84
1	185	5	Net, The	Action Crime Thriller	1996-08-02 10:58:45	1995	1996-08-04	3.538801	3.129334	13469	23	586
1	292	5	Outbreak	Action Drama Sci-Fi Thriller	1996-08-02 10:57:01	1995	1996-08-04	3.538801	3.418011	14447	23	628
1	316	5	Stargate	Action Adventure Sci-Fi	1996-08-02 10:56:32	1994	1996-08-04	3.538801	3.349677	17030	24	710
1	329	5	Star Trek: Generations	Action Adventure Drama Sci-Fi	1996-08-02 10:56:32	1994	1996-08-04	3.538801	3.337457	14550	24	606

```
tibble [9,000,055 x 13] (S3: tbl_df/tbl/data.frame)
$ userId      : int [1:9000055] 1 1 1 1 1 1 1 1 1 1 ...
$ movieId     : num [1:9000055] 122 185 292 316 329 355 356 362 364 370 ...
$ rating      : num [1:9000055] 5 5 5 5 5 5 5 5 5 5 ...
$ title       : chr [1:9000055] "Boomerang" "Net, The" "Outbreak" "Stargate" ...
$ genres      : chr [1:9000055] "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|Sci-Fi|Thriller"
$ rating_date : POSIXct[1:9000055], format: "1996-08-02 11:24:06" "1996-08-02 10:58:45" ...
$ movie_dt    : num [1:9000055] 1992 1995 1995 1994 1994 ...
```

```
$ date          : POSIXct[1:9000055], format: "1996-08-04" "1996-08-04" ...
$ avg_rating    : num [1:9000055] 3.54 3.54 3.54 3.54 3.54 ...
$ avg_rating_rel: num [1:9000055] 2.86 3.13 3.42 3.35 3.34 ...
$ n            : int [1:9000055] 2178 13469 14447 17030 14550 4831 31079 3612 18921 7331 ...
$ years        : num [1:9000055] 26 23 23 24 24 24 24 24 24 24 ...
$ n_year       : num [1:9000055] 84 586 628 710 606 ...
```

Modify validation data for Release Date Effect We need to do the above `avg_rating_rel_effect` update for the validation data as well. Let's update the `validation` table to get Table 11 below:

All Years

Table 11: MovieLens validation data with average rating due to release date effect

userId	movieId	rating	title	genres	rating_date	movie_dt	date	avg_rating	avg_rating_rel	n	years	n_year
1	231	5	Dumb & Dumber	Comedy	1996-08-02 10:56:32	1994	1996-08-04	3.555820	2.953281	1798	24	75
1	480	5	Jurassic Park	Action Adventure Sci-Fi Thriller	1996-08-02 11:00:53	1993	1996-08-04	3.555820	3.643993	3271	25	131
1	586	5	Home Alone	Children Comedy	1996-08-02 11:07:48	1990	1996-08-04	3.555820	3.074550	1556	28	56
2	151	3	Rob Roy	Action Drama Romance War	1997-07-07 03:34:10	1995	1997-07-06	3.606571	3.571984	771	23	34
2	858	2	Godfather, The	Crime Drama	1997-07-07 03:20:45	1972	1997-07-06	3.606571	4.412675	2067	46	45

```
tibble [999,999 x 13] (S3: tbl_df/tbl/data.frame)
$ userId      : int [1:999999] 1 1 1 2 2 2 3 3 4 4 ...
$ movieId     : num [1:999999] 231 480 586 151 858 ...
$ rating      : num [1:999999] 5 5 5 3 2 3 3.5 4.5 5 3 ...
$ title       : chr [1:999999] "Dumb & Dumber" "Jurassic Park" "Home Alone" "Rob Roy" ...
$ genres      : chr [1:999999] "Comedy" "Action|Adventure|Sci-Fi|Thriller" "Children|Comedy" "Action"
$ rating_date : POSIXct[1:999999], format: "1996-08-02 10:56:32" "1996-08-02 11:00:53" ...
$ movie_dt    : num [1:999999] 1994 1993 1990 1995 1972 ...
$ date        : POSIXct[1:999999], format: "1996-08-04" "1996-08-04" ...
$ avg_rating   : num [1:999999] 3.56 3.56 3.56 3.61 3.61 ...
$ avg_rating_rel: num [1:999999] 2.95 3.64 3.07 3.57 4.41 ...
$ n           : int [1:999999] 1798 3271 1556 771 2067 862 2545 947 1869 776 ...
$ years       : num [1:999999] 24 25 28 23 46 21 28 17 23 24 ...
$ n_year      : num [1:999999] 75 131 56 34 45 41 91 56 81 32 ...
```

3 Analysis - Model Building and Evaluation

3.1 Split the edx data into separate training and test sets

We will develop our algorithm using the edx set only.

We will split the edx data into separate training and test sets to design and test our algorithm, namely train_set and test_set.

3.1.1 Loss function

For a final test of our algorithm, we predict movie ratings in the test set as if they were unknown. RMSE⁷ (residual mean squared error/root mean square error), the typical error loss, will be used to evaluate how close our predictions are to the true values in the validation set.

We define $y_{u,i}$ as the rating for movie i by user u and denote our prediction with $\hat{y}_{u,i}$.

The RMSE is then defined as Equation 2:

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2} \quad (2)$$

with N being the number of user/movie combinations and the sum occurring over all these combinations.

Remember that we can interpret the RMSE similarly to a standard deviation: it is the typical error we make when predicting a movie rating. If this number is larger than 1, it means our typical error is larger than one star, which is not good.

Let's write a function that computes the RMSE for vectors of ratings and their corresponding predictors:

```
RMSE <- function(true_ratings, predicted_ratings) {  
  sqrt(mean((true_ratings - predicted_ratings)^2))  
}
```

⁷https://en.wikipedia.org/wiki/Root-mean-square_deviation

3.2 Model 1: A first naive “mean” model

Let's start by building the simplest possible recommendation system: we predict the same rating for all movies regardless of user. A model that assumes the same rating for all movies and users with all the differences explained by random variation would look like Equation 3:

$$Y_{i,i} = \mu + \epsilon_{u,i} \quad (3)$$

with $\epsilon_{u,i}$ independent errors sampled from the same distribution centered at 0 and μ the “true” rating for all movies. We know that the estimate that minimizes the RMSE is the least squares estimate of μ and, in this case, is the average of all ratings:

```
(mu_hat <- mean(train_set$rating))
[1] 3.512482
```

If we predict all unknown ratings with $\hat{\mu}$ we obtain the following RMSE:

```
(model_1_rmse <- RMSE(test_set$rating, mu_hat))
[1] 1.059904
```

Keep in mind that if we plug in any other number, we get a higher RMSE. For example:

```
predictions <- rep(2.5, nrow(test_set))
RMSE(test_set$rating, predictions)
[1] 1.465736

predictions <- rep(3, nrow(test_set))
RMSE(test_set$rating, predictions)
[1] 1.177271

predictions <- rep(4, nrow(test_set))
RMSE(test_set$rating, predictions)
[1] 1.166678
```

From looking at the distribution of ratings, we can visualize that this is the standard deviation of that distribution. We get a RMSE of about 1. Our target is RMSE < 0.86490. So we can definitely do better!

3.2.1 Results Model 1

As we go along, we will be comparing different approaches. Let's start by creating a results table with this naive approach to get Table 12:

Table 12: RMSE Results Model 1

Index	Method	RMSE
1	Just the average	1.059904

3.3 Model 2: Movie effects

We know from experience that some movies are just generally rated higher than others. This intuition, that different movies are rated differently, is confirmed by data. We can augment our previous model by adding the term b_i to represent average ranking for movie i and would look like Equation 4:

$$Y_{u,i} = \mu + b_i + \epsilon_{u,i} \quad (4)$$

Statistics textbooks refer to the b s as effects or “bias”.

We can again use least squares to estimate the b_i in the following way to get Equation 5:

$$fit \leftarrow lm(rating \sim as.factor(userId), data = train_set) \quad (5)$$

Because there are thousands of b_i as each movie gets one, the `lm()` function will be very slow here. We therefore will not run the code above.

But in this particular situation, we know that the least squares estimate \hat{b}_i is just the average of $Y_{u,i} - \hat{\mu}$ for each movie i . *So we can compute them this way (we will drop the hat notation in the code to represent estimates going forward):*

$$\hat{b}_i = \overline{y_{u,i} - \hat{\mu}} \quad (6)$$

```
mu <- mean(train_set$rating)
```

We can see that these estimates vary substantially, see Figure 11

Movie effect or bias distribution

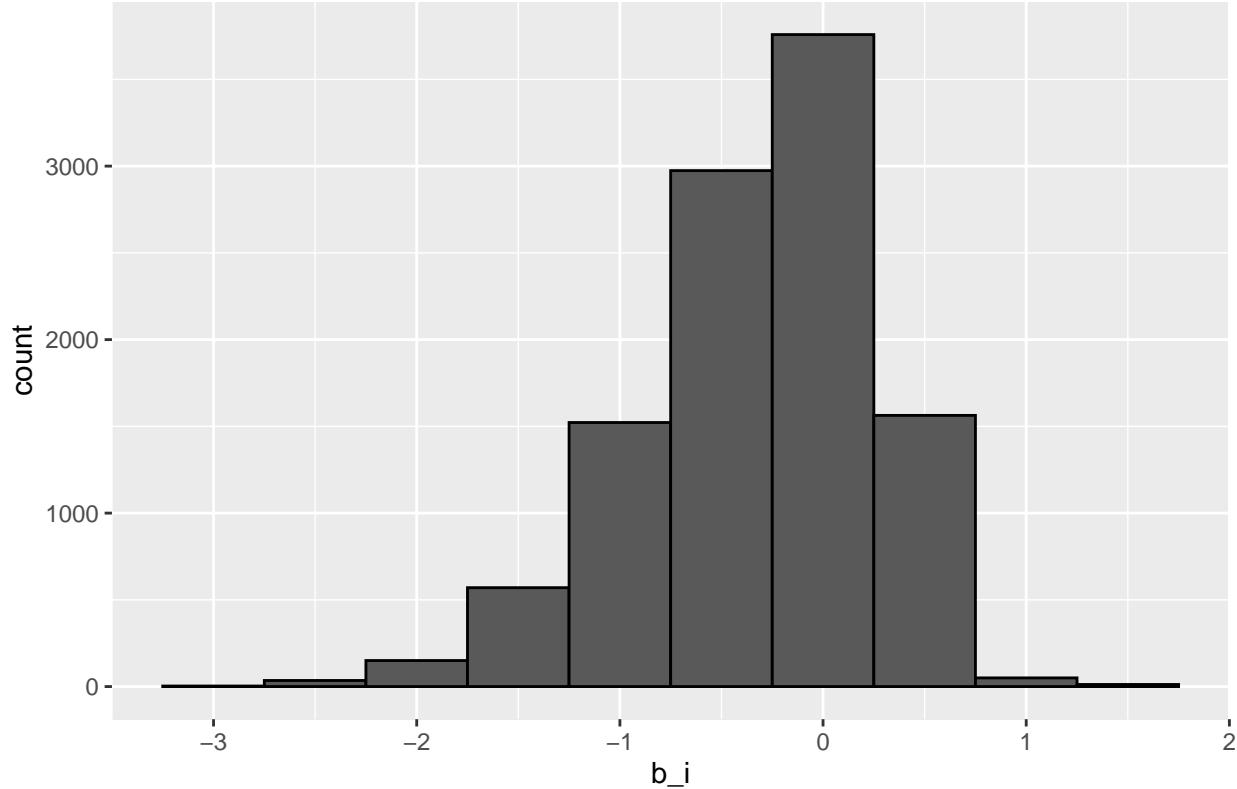


Figure 11: Movie effect or bias distribution

Remember $\hat{\mu}=3.5$ so a $b_i=1.5$ implies a perfect five star rating. Let's see how much our prediction improves once we use $y_{u,i}^* = \hat{\mu} + \hat{b}_i$:

```
predicted_ratings_model_2 <- mu + test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  .$b_i
```

```
[1] 0.9437429
```

3.3.1 Results Model 1-2

Let's add the movie effects model to our results table to get Table 13

Table 13: RMSE Results Models 1-2

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429

3.4 Model 3: User effects

Let's compute b_u the average rating for user u for those that have rated over 100 movies, see Figure 12

```
train_set %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating)) %>%
  filter(n()>=100) %>%
  ggplot(aes(b_u)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for users who have rated over 100 movies")
```

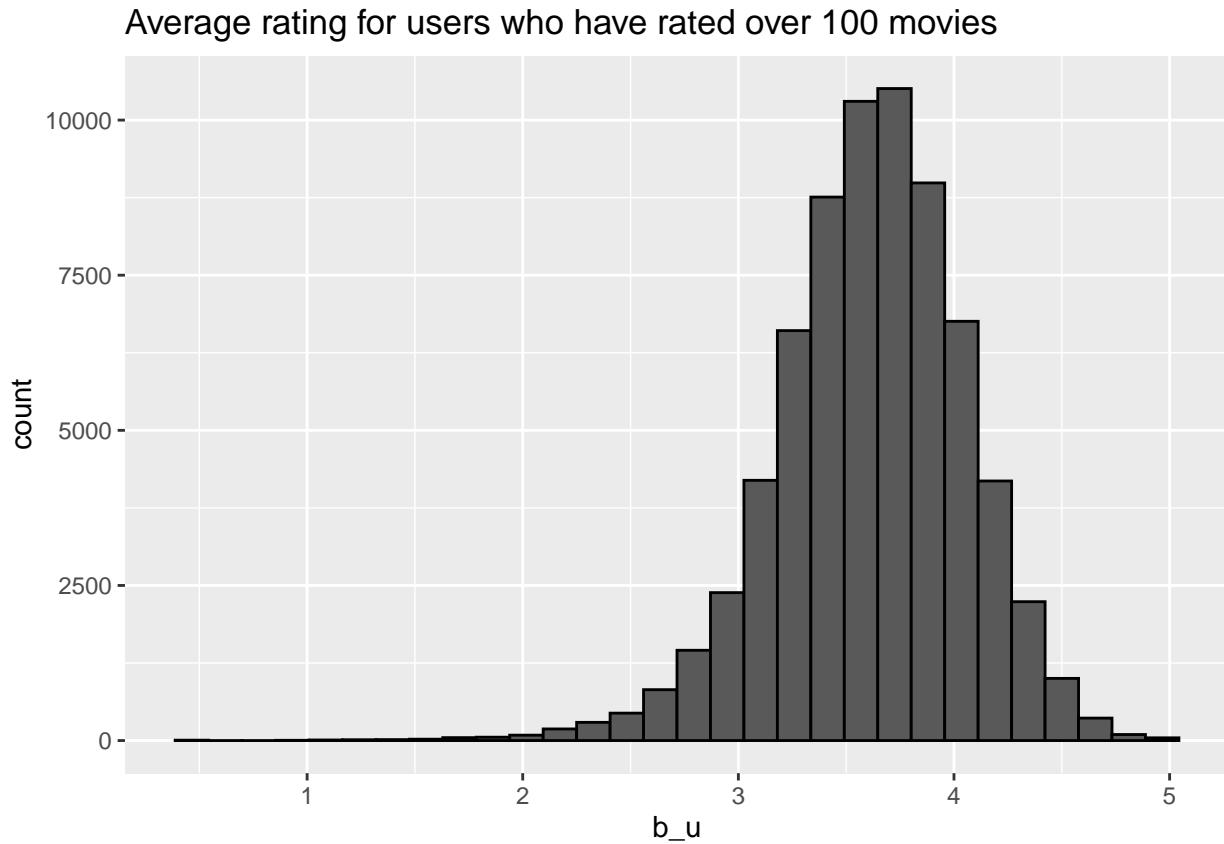


Figure 12: Average rating for users who have rated over 100 movies

Let's compute b_u the average rating for user u for those that have rated any movies, see Figure 13

```
train_set %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating)) %>%
  ggplot(aes(b_u)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for users who have rated any movies")
```

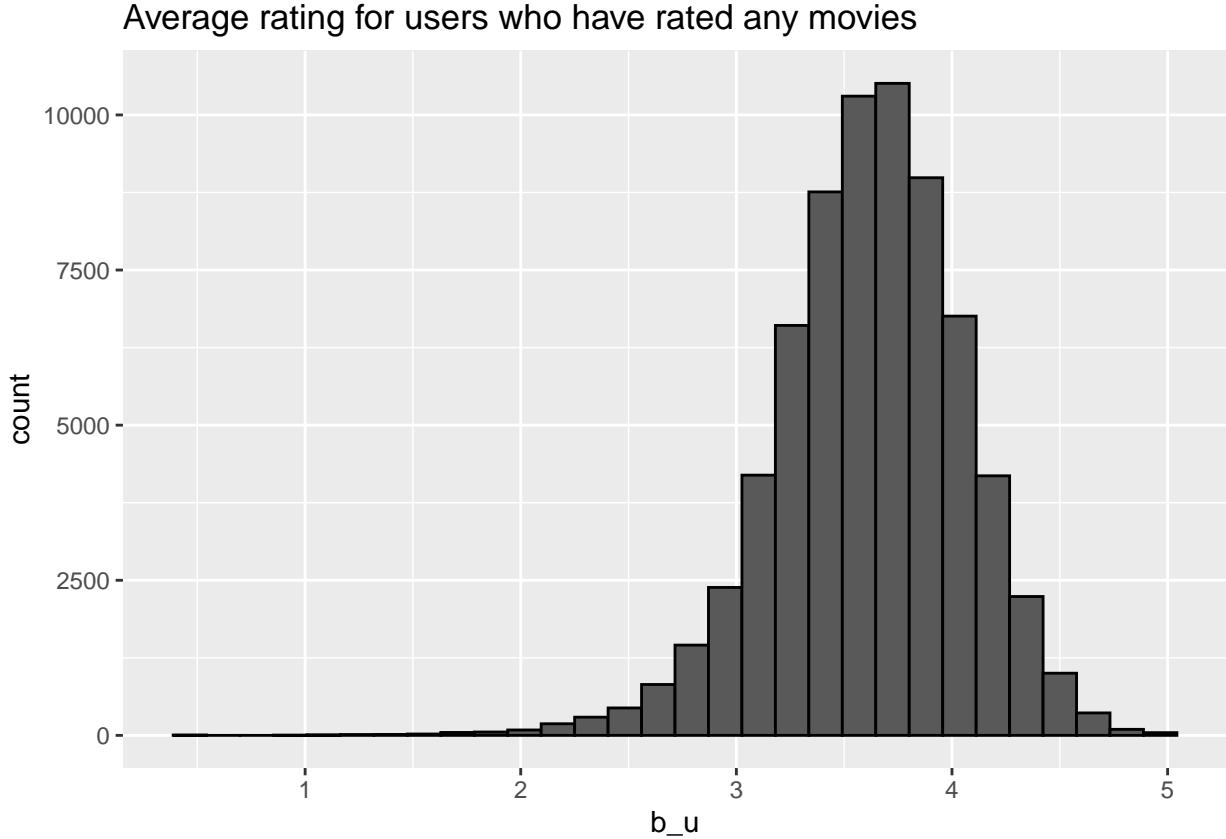


Figure 13: Average rating for users who have rated any movies

Notice that there is substantial variability across users as well: some users are very cranky and others love every movie. This implies that a further improvement to our model may be as shown in Equation 7:

$$Y_{u,i} = \mu + b_i + b_u + \epsilon_{u,i} \quad (7)$$

where b_u is a user-specific effect. Now if a cranky user (negative b_u) rates a great movie (positive b_i), the effects counter each other and we may be able to correctly predict that this user gave this great movie a 3 rather than a 5.

To fit this model, we could again use `lm()` as shown in Equation 8:

$$fit \leftarrow lm(rating \sim as.factor(movieId) + as.factor(userId), data = train_set) \quad (8)$$

but, for the reasons described earlier, we won't. Instead, we will compute an approximation by computing $\hat{\mu}$ and \hat{b}_i and estimating \hat{b}_u as the average of $y_{u,i} - \hat{\mu} - \hat{b}_i$:

$$\hat{b_u} = \overline{y_{u,i} - \hat{\mu} - \hat{b}_i} \quad (9)$$

```

user_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating - mu - b_i))

```

We can see that these estimates vary substantially, see Figure 14

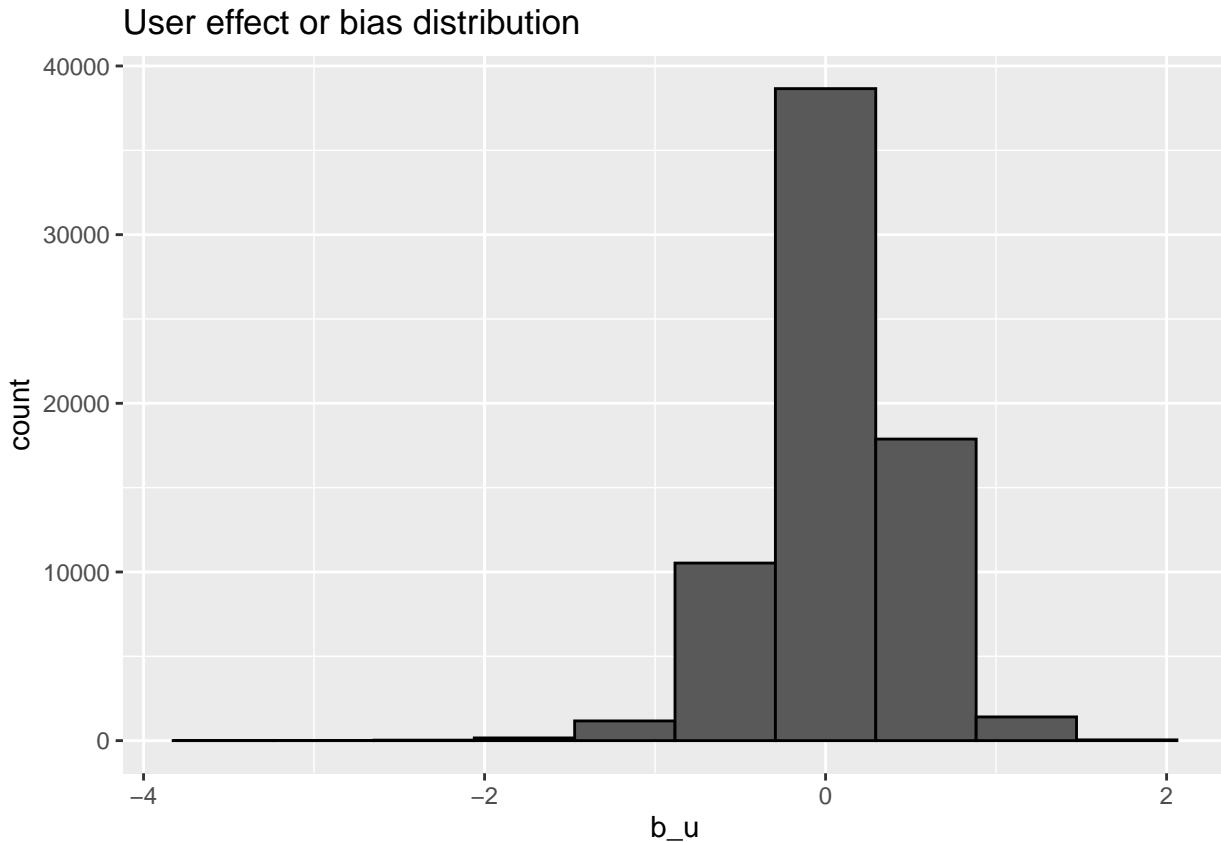


Figure 14: User effect or bias distribution

We can now construct predictors and see how much the RMSE improves:

```

predicted_ratings_model_3 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>%
  .$pred

```

[1] 0.865932

3.4.1 Results Table Model 1-3

Let's add the user effects model to our results table to get Table 14

Table 14: RMSE Results Models 1-3

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
3	Movie + User Effects Model	0.8659320

3.5 Model 4: Genre effects

The movielens data also has a genres column. This column includes every genre that applies to the movie. Some movies fall under several genres. Define a category of genres as whatever combination of genres appears in this column. We will refer to this category as simply “genres”.

There is strong evidence of a genre effect as we have shown earlier in Figure 4, and in this section below in Figure 16. If we define $g_{u,i}$ as the genre for u user’s rating of movie i , then the following model as shown in Equation 10 is most appropriate:

$$Y_{u,i} = \mu + b_i + b_u + \sum_{k=1}^K x_{u,i}\beta_k + \epsilon_{u,i} \quad (10)$$

with $x_{u,i}^k = 1$ if $g_{u,i}$ is genre *k*

```
train_set %>%
  group_by(genres) %>%
  summarize(mu_g = mean(rating)) %>%
  ggplot(aes(mu_g)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for movies of category genres")
```

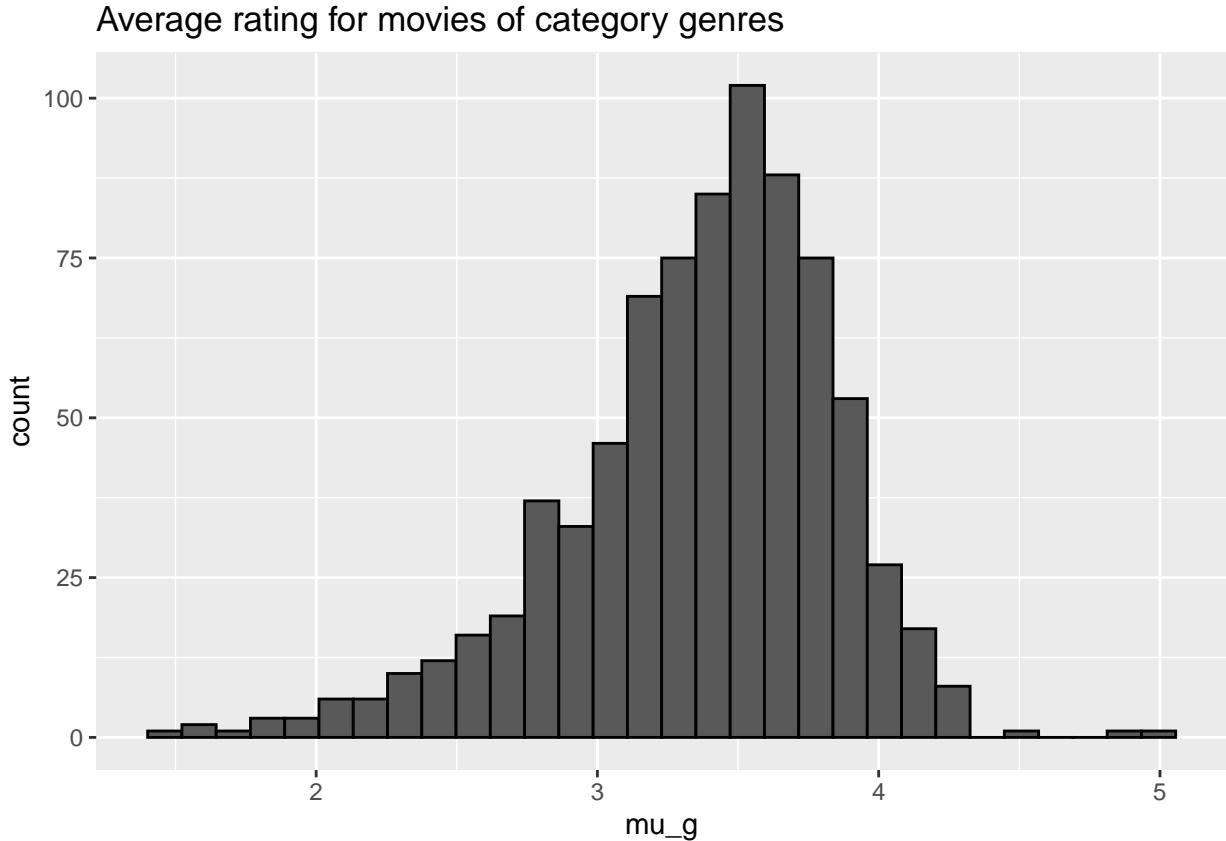


Figure 15: Average rating for movies of category genres

To fit this model, we could again use the `lm()` function as shown in Equation 11:

$$fit \leftarrow lm(rating \sim as.factor(movieId) + as.factor(userId) + as.factor(genres), data = train_set) \quad (11)$$

but, for the reasons described earlier, we won't. Instead, we will compute an approximation by computing $\hat{\mu}$, \hat{b}_i , \hat{b}_u and estimating \hat{b}_g as the average of $y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u$:

$$\hat{b}_g = \overline{y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u} \quad (12)$$

where:

$$\hat{b}_g = \sum_{k=1}^K x_{u,i} \beta_k \quad (13)$$

with $x_{u,i}^k = 1$ if $g_{u,i}$ is genre $*k*$

where \hat{b}_g is genre specific effect.

```
genres_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = mean(rating - mu - b_i - b_u))
```

We can see that these estimates vary substantially, see Figure 16

We can now construct predictors and see how much the RMSE improves:

```
predicted_ratings_model_4 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  mutate(pred = mu + b_i + b_u + b_g) %>%
  .$pred
```

[1] 0.8655941

Genres effect or bias distribution

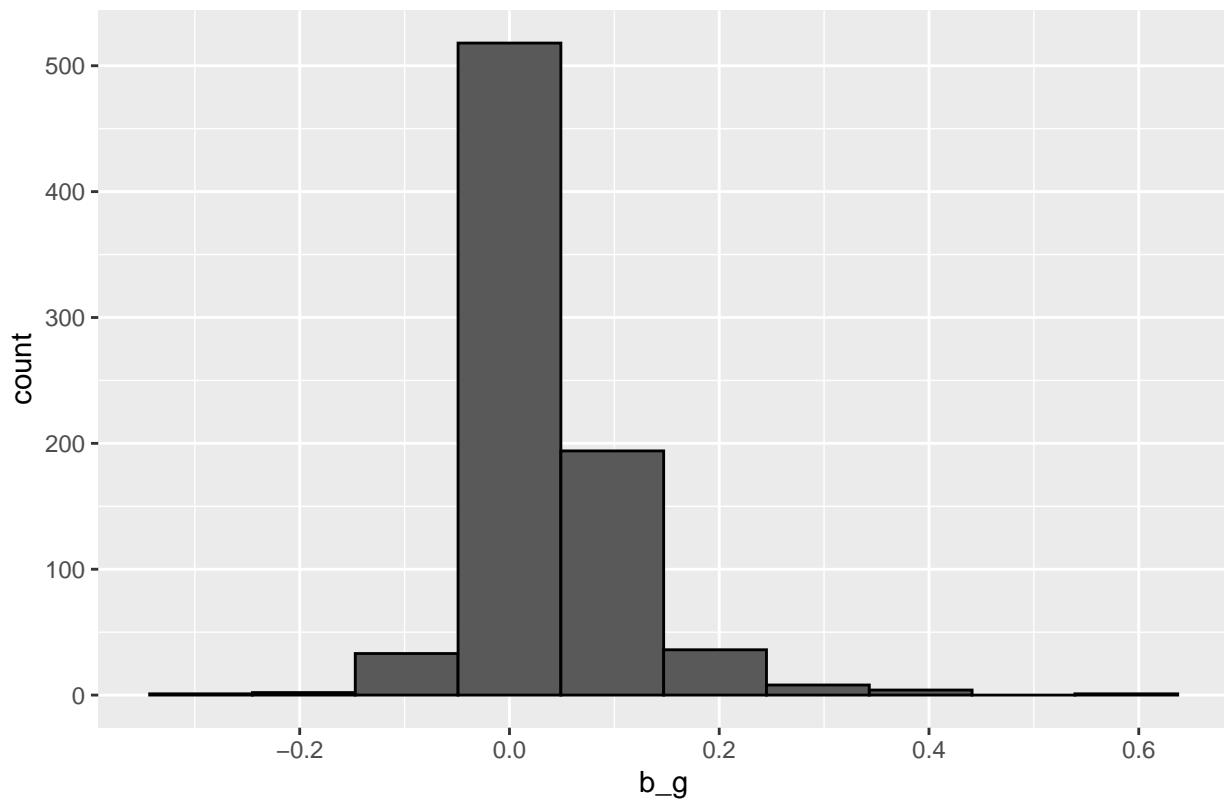


Figure 16: Genres effect or bias distribution

3.5.1 Results Table Model 1-4

Let's add the genres effects model to our results table to get Table 15

Table 15: RMSE Results Models 1-4

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941

3.6 Model 5: Rating Time effect

The movielens dataset also includes a time stamp. This variable represents the time and date in which the rating was provided. Earlier in the EDA/Data wrangling section we created a new column date with the time stamp.

We computed the average rating for each week and plotted this average against day.

The plot shows some evidence of a time effect. If we define $d_{u,i}$ as the day for user's u rating of movie i , then the following updated model as shown in Equation 14 is most appropriate:

$$Y_{u,i} = \mu + b_i + b_u + \sum_{k=1}^K x_{u,i}\beta_k + f(d_{u,i}) + \epsilon_{u,i} \quad (14)$$

with f a smooth function of $d_{u,i}$

To fit this model, we could again use `lm()` function as shown in Equation 15:

$$\begin{aligned} \text{fit} \leftarrow \text{lm}(rating \sim \text{as.factor(movieId)} + \text{as.factor(userId)} + \\ \text{as.factor(genres)} + \text{as.factor(date)}, \text{data} = \text{train_set}) \end{aligned} \quad (15)$$

but, for the reasons described earlier, we won't. Instead, we will compute an approximation by computing $\hat{\mu}$, \hat{b}_i , \hat{b}_u , \hat{b}_g and estimating \hat{b}_d as the average of $y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u - \hat{b}_g$:

$$\hat{b}_d = \overline{y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u - \hat{b}_g} \quad (16)$$

where:

$$\hat{b}_d = f(d_{u,i}) \quad (17)$$

where \hat{b}_d is rating time specific effect.

```
time_effect_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  group_by(date) %>%
  summarize(b_d = mean(avg_rating - mu - b_i - b_u - b_g))
```

We can see that these estimates vary substantially, see Figure 16

Rating time effect or bias distribution

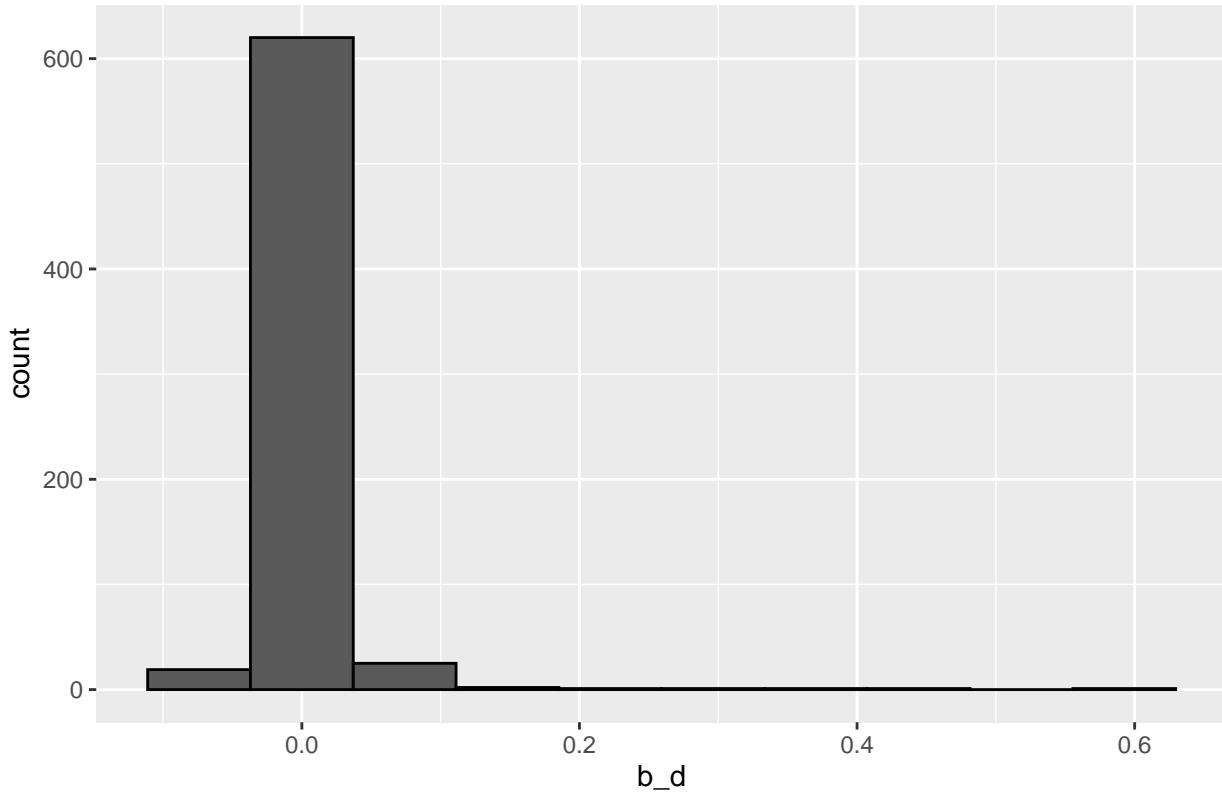


Figure 17: Rating time effect or bias distribution

We can now construct predictors and see how much the RMSE improves

```
predicted_ratings_model_5 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  left_join(time_effect_avgs, by = "date") %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d) %>%
  .$pred
```

[1] 0.8654205

3.6.1 Results Table Model 1-5

Let's add the Rating Time effects model to our results table to get Table 16

Table 16: RMSE Results Models 1-5

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205

3.7 Model 6: Release Date Effect

The plots in Figures 6, 7, 8, 9, 10 above shows some evidence of a Release Date effect based on the when the movie was released and it's popularity given by the mean rating. If we define $arr_{r,i,y}$ as the average rating $r=mean(rating)$ since release date $y=n_year$ for movie i (in the formula for plots above), then the following updated model is most appropriate:

$$Y_{u,i} = \mu + b_i + b_u + \sum_{k=1}^K x_{u,i}\beta_k + f(d_{u,i}) + f(arr_{r,i,y}) + \epsilon_{u,i} \quad (18)$$

with f a smooth function of $arr_{r,i,y}$

To fit this model, we could again use `lm()` function as shown in Equation 19:

$$\begin{aligned} fit \leftarrow lm(rating \sim as.factor(movieId) + as.factor(userId) + \\ as.factor(genres) + as.factor(date) + \\ as.factor(movie_dt), data = train_set) \end{aligned} \quad (19)$$

but, for the reasons described earlier, we won't. Instead, we will compute an approximation by computing $\hat{\mu}, \hat{b}_i, \hat{b}_u, \hat{b}_g, \hat{b}_d$ and estimating \hat{b}_r as the average of $y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u - \hat{b}_g - \hat{b}_d$ where:

$$\hat{b}_r = \overline{y_{u,i} - \hat{\mu} - \hat{b}_i - \hat{b}_u - \hat{b}_g - \hat{b}_d} \quad (20)$$

where:

$$\hat{b}_r = f(arr_{r,i,y}) \quad (21)$$

where \hat{b}_r is Release date specific effect.

```
rel_effect_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  left_join(time_effect_avgs, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = mean(avg_rating_rel - mu - b_i - b_u - b_g - b_d))
```

We can see that these estimates vary substantially, see Figure 18

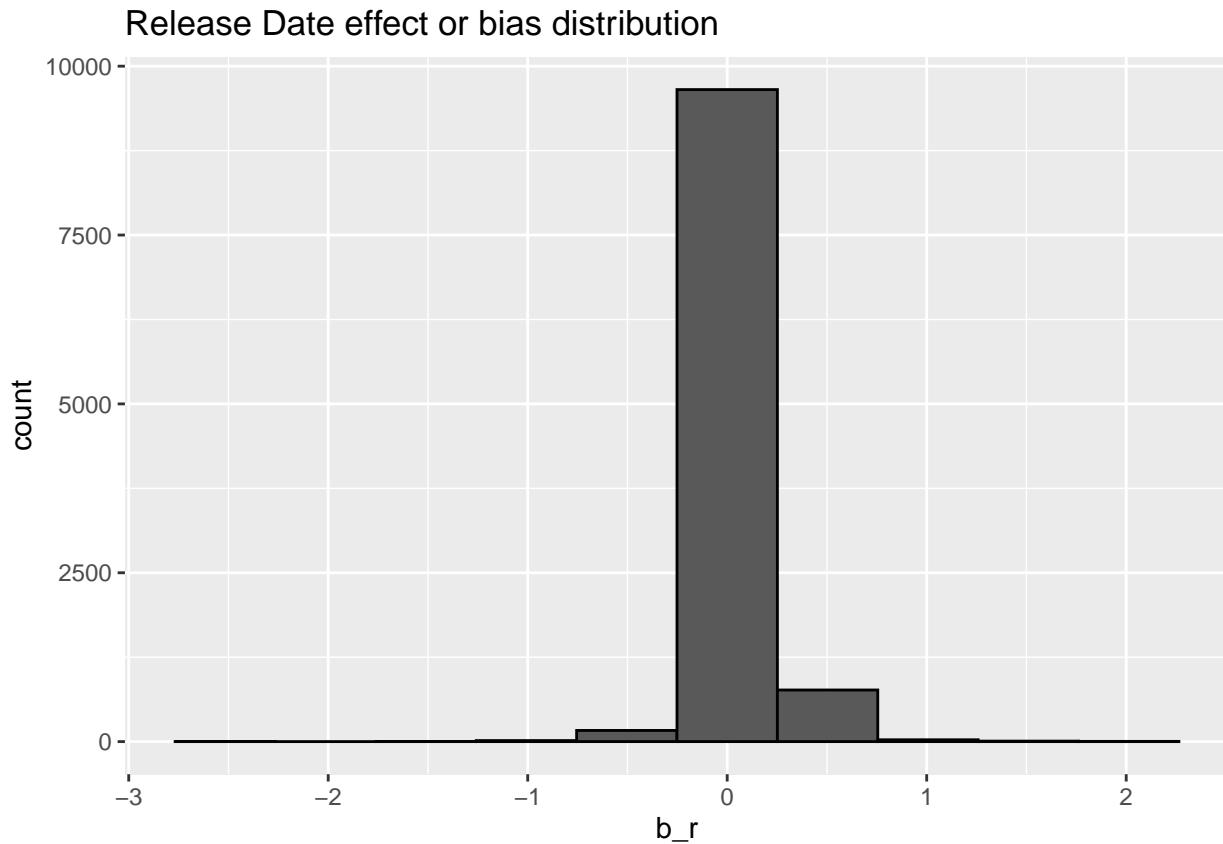


Figure 18: Release Date effect or bias distribution

We can now construct predictors and see how much the RMSE improves

```
predicted_ratings_model_6 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  left_join(time_effect_avgs, by = "date") %>%
  left_join(rel_effect_avgs, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .$pred

(model_6_rmse <- RMSE(predicted_ratings_model_6, test_set$rating))

[1] 0.863333
```

3.7.1 Results Table Model 1-6

Let's add the Release Date effects model to our results table to get Table 17

Table 17: RMSE Results Models 1-6

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330

3.8 Regularization

3.8.1 Motivation

Despite the large movie to movie variation, our improvement in RMSE is either relatively negligible or the results from the recommendations are strange. Let's explore where we made mistakes in our second model, using only movie effects b_i .

Here are the **10 largest mistakes**

```
test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  mutate(residual = rating - (mu + b_i)) %>%
  arrange(desc(abs(residual))) %>%
  slice(1:10) %>%
  pull(title) %>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="",
        caption="Without Regularization 10 largest mistakes\\label{tbl:without_regularization_10_largest_mistakes}",
        kable_styling(latex_options=c("HOLD_position")))
```

Table 18: Without Regularization 10 largest mistakes

x
From Justin to Kelly
Time Changer
The Shawshank Redemption, The
Children Underground

These all seem like obscure movies, or in this case a repetition. Many of them have large predictions. Let's look at the **top 10 worst and best movies** based on \hat{b}_i . First, let's create a database that connects 'movieId' to movie title

```
movie_titles <- edx %>%
  select(movieId, title) %>%
  distinct()
```

Here are the **10 best movies** according to our estimate

```
movie_avgs %>% left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  slice(1:10) %>%
  pull(title)%>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Without Regularization 10 Best movies",
        kable_styling(latex_options=c("HOLD_position")))
```

Table 19: Without Regularization 10 Best movies

x
Hellhounds on My Trail (1999)
Who's Singin' Over There? (a.k.a. Who Sings Over There) (Ko to tamo peva) (1980)
Satan's Tango (SĀjtĀntangĀ³) (1994)
Shadows of Forgotten Ancestors (1964)
Money (Argent, L') (1983)
Fighting Elegy (Kenka erejii) (1966)
Sun Alley (Sonnenallee) (1999)
Aerial, The (La Antena) (2007)
Blue Light, The (Das Blaue Licht) (1932)
More (1998)

And here are the **10 worst movies**

```
movie_avgs %>% left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %>%
  slice(1:10)  %>%
  pull(title)%>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Without Regularization 10 Worst movies")
  kable_styling(latex_options=c("HOLD_position"))
```

Table 20: Without Regularization 10 Worst movies

x
Besotted (2001)
Confessions of a Superhero (2007)
War of the Worlds 2: The Next Wave (2008)
SuperBabies: Baby Geniuses 2 (2004)
From Justin to Kelly (2003)
Legion of the Dead (2000)
Disaster Movie (2008)
Hip Hop Witch, Da (2000)
Criminals (1996)
Mountain Eagle, The (1926)

They all seem to be quite obscure. Let's look at how often they are rated.

10 best movies

```
train_set %>% count(movieId) %>%
  left_join(movie_avgs, by="movieId") %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  slice(1:10) %>%
  pull(n)
[1] 1 3 2 1 1 1 1 1 1 6
```

10 worst movies

```

train_set %>% count(movieId) %>%
  left_join(movie_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %>%
  slice(1:10) %>%
  pull(n)
[1] 1 1 2 40 168 4 28 11 1 1

```

The *supposed “best” and “worst” movies were rated by very few users*, in most cases just 1. These movies were mostly obscure ones. This is because with just a few users, we have more uncertainty. Therefore, larger estimates of b_i , negative or positive, are more likely.

These are noisy estimates that we should not trust, especially when it comes to prediction. Large errors can increase our RMSE, so we would rather be conservative when unsure.

Regularization permits us to penalize large estimates that are formed using small sample sizes.

3.8.2 Penalized least squares

The general idea behind regularization is to constrain the total variability of the effect sizes. Why does this help? Consider a case in which we have movie $i = 1$ with 100 user ratings and 4 movies $i = 2, 3, 4, 5$ with just one user rating. We intend to fit the model

$$Y_{u,i} = \mu + b_i + \epsilon_{u,i}$$

Suppose we know the average rating is, say, $\mu = 3$. If we use least squares, the estimate for the first movie effect b_1 is the average of the 100 user ratings, $\frac{1}{100} \sum_{i=1}^{100} (Y_{i,1} - \mu)$, which we expect to be a quite precise. However, the estimate for movies 2, 3, 4, and 5 will simply be the observed deviation from the average rating $\hat{b}_i = Y_{u,i} - \hat{\mu}$ which is an estimate based on just one number so it won't be precise at all. Note these estimates make the error $Y_{u,i} - \mu + \hat{b}_i$ equal to 0 for $i = 2, 3, 4, 5$, but this is a case of over-training. In fact, ignoring the one user and guessing that movies 2, 3, 4, and 5 are just average movies ($b_i = 0$) might provide a better prediction. The general idea of penalized regression is to control the total variability of the movie effects: $\sum_{i=1}^5 b_i^2$. Specifically, instead of minimizing the least squares equation, we minimize an equation that adds a penalty:

$$\frac{1}{N} \sum_{u,i} (y_{u,i} - \mu - b_i)^2 + \lambda \sum_i b_i^2$$

The first term is just least squares and the second is a penalty that gets larger when many b_i are large. Using calculus we can actually show that the values of b_i that minimize this equation are:

$$\hat{b}_i(\lambda) = \frac{1}{\lambda + n_i} \sum_{u=1}^{n_i} (Y_{u,i} - \hat{\mu})$$

where n_i is the number of ratings made for movie i . This approach will have our desired effect: when our sample size n_i is very large, a case which will give us a stable estimate, then the penalty λ is effectively ignored since $n_i + \lambda \approx n_i$. However, when the n_i is small, then the estimate $\hat{b}_i(\lambda)$ is shrunken towards 0. The larger λ , the more we shrink.

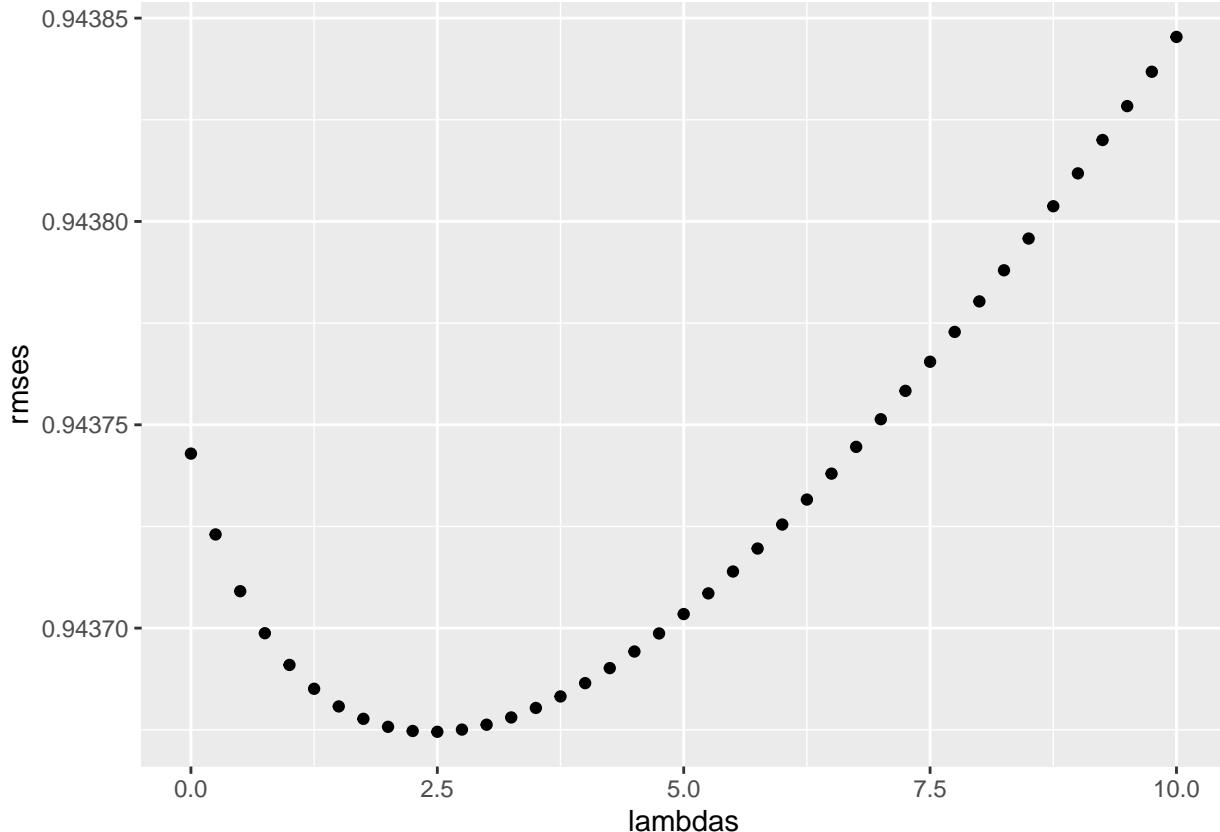
3.9 Model 7 : Model 2 + Regularization: Choosing the penalty terms

Note that λ is a tuning parameter. We will use cross-validation to choose it.

Here we use only 1 fold

```
lambdas <- seq(0, 10, 0.25)
mu <- mean(train_set$rating)
just_the_sum <- train_set %>%
  group_by(movieId) %>%
  summarize(s = sum(rating - mu), n_i = n())
rmses <- sapply(lambdas, function(l){
  predicted_ratings <- test_set %>%
    left_join(just_the_sum, by='movieId') %>%
    mutate(b_i = s/(n_i+1)) %>%
    mutate(pred = mu + b_i) %>%
    .$pred
  return(RMSE(predicted_ratings, test_set$rating))
})
(lambda <- lambdas[which.min(rmses)])
[1] 2.5
```

```
qplot(lambdas, rmses)
```



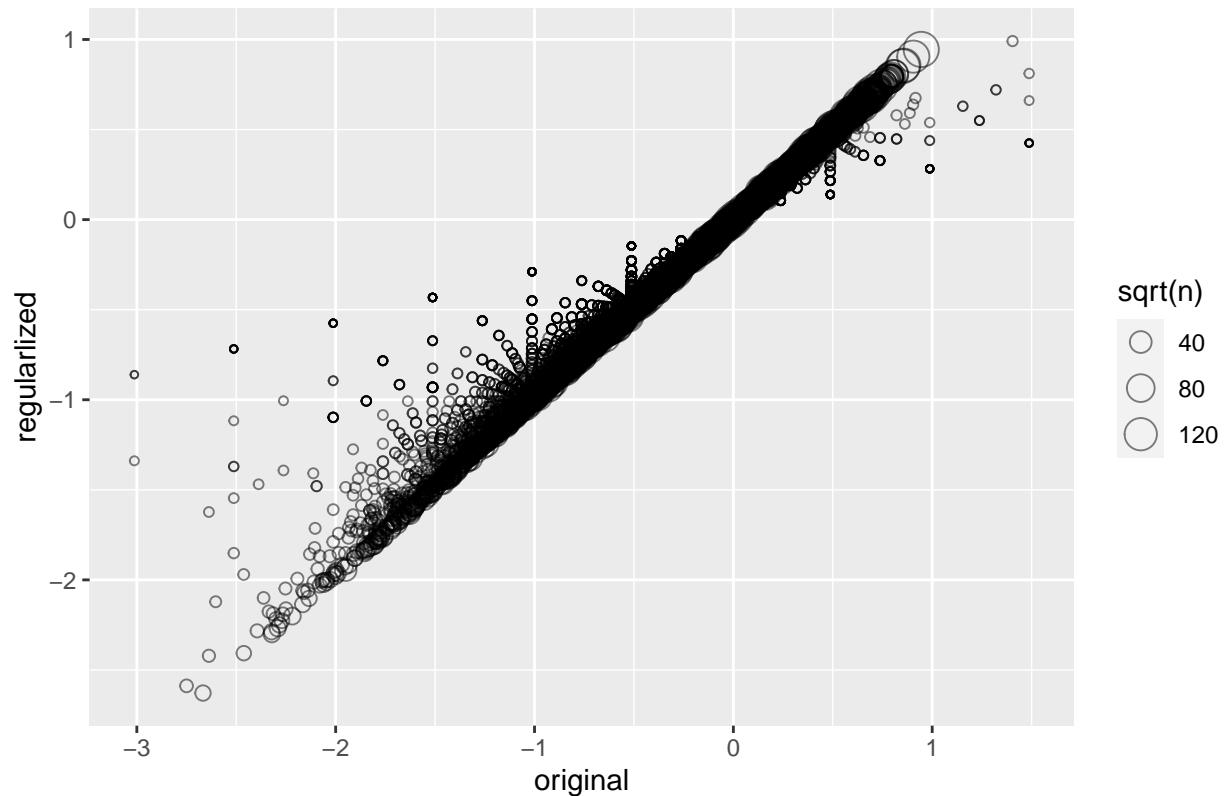
```

mu <- mean(train_set$rating)
movie_reg_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n() + lambda), n_i = n())

```

To see how the estimates shrink, let's make a plot of the regularized estimates versus the least squares estimates.

Movie Effect Reg



3.9.1 Effect of use of penalized estimates

Now, let's look at the top 10 best movies based on the penalized estimates $\hat{b}_i(\lambda)$

Top 10 best movies

```
train_set %>%
  dplyr::count(movieId) %>%
  left_join(movie_reg_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  select(title, b_i, n) %>%
  slice(1:10) %>%
  kable("latex", escape=TRUE, booktabs=TRUE, linesep="", caption="With Regularization top 10 best movies
Joining, by = "movieId"
```

Table 21: With Regularization top 10 best movies based on the penalized estimates

title	b_i	n
More (1998)	0.9911891	6
Shawshank Redemption, The (1994)	0.9447301	22363
Godfather, The (1972)	0.9048453	14107
Usual Suspects, The (1995)	0.8568137	17315
Schindler's List (1993)	0.8514631	18567
Who's Singin' Over There? (a.k.a. Who Sings Over There) (Ko to tamo peva) (1980)	0.8113734	3
Rear Window (1954)	0.8087002	6325
Casablanca (1942)	0.8046208	9027
Double Indemnity (1944)	0.7972822	1711
Third Man, The (1949)	0.7961995	2420

These make much more sense! These movies are watched more and have more ratings.

Top 10 worst movies

```

train_set %>%
  dplyr::count(movieId) %>%
  left_join(movie_reg_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %>%
  select(title, b_i, n) %>%
  slice(1:10) %>%
  kable("latex", escape=TRUE, booktabs=TRUE, linesep="", caption="With Regularization top 10 worst movies")
  kable_styling(latex_options=c("HOLD_position"))
Joining, by = "movieId"

```

Table 22: With Regularization top 10 worst movies based on the penalized estimates

title	b_i	n
From Justin to Kelly (2003)	-2.628135	168
SuperBabies: Baby Geniuses 2 (2004)	-2.588218	40
Disaster Movie (2008)	-2.421295	28
PokÃ©mon Heroes (2003)	-2.406821	109
Glitter (2001)	-2.300597	282
Barney's Great Adventure (1998)	-2.287959	168
Carnosaur 3: Primal Species (1996)	-2.282927	51
Gigli (2003)	-2.268819	249
Pokemon 4 Ever (a.k.a. PokÃ©mon 4: The Movie) (2002)	-2.249836	168
Son of the Mask (2005)	-2.226802	128

Improved our results

```

predicted_ratings_model_7 <- test_set %>%
  left_join(movie_reg_avgs, by='movieId') %>%
  mutate(pred = mu + b_i) %>%
  .$pred

```

[1] 0.9436745

3.9.2 Results Table Model 1-7

Let's add Regularization to our Model 1-2 results table to get Table 23

Table 23: RMSE Results Models 1-7

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330
7	Regularized Movie Effect Model - 1 fold CV	0.9436745

The penalized estimates provide an improvement over the least squares estimates.

3.10 Model 8 : Model 4 + Regularization: Choosing the penalty terms

We will be using multi-fold full cross-validation on the entire edx data set shortly.

But first... we use regularization for the estimate user effects as well as the genres effects. We are minimizing:

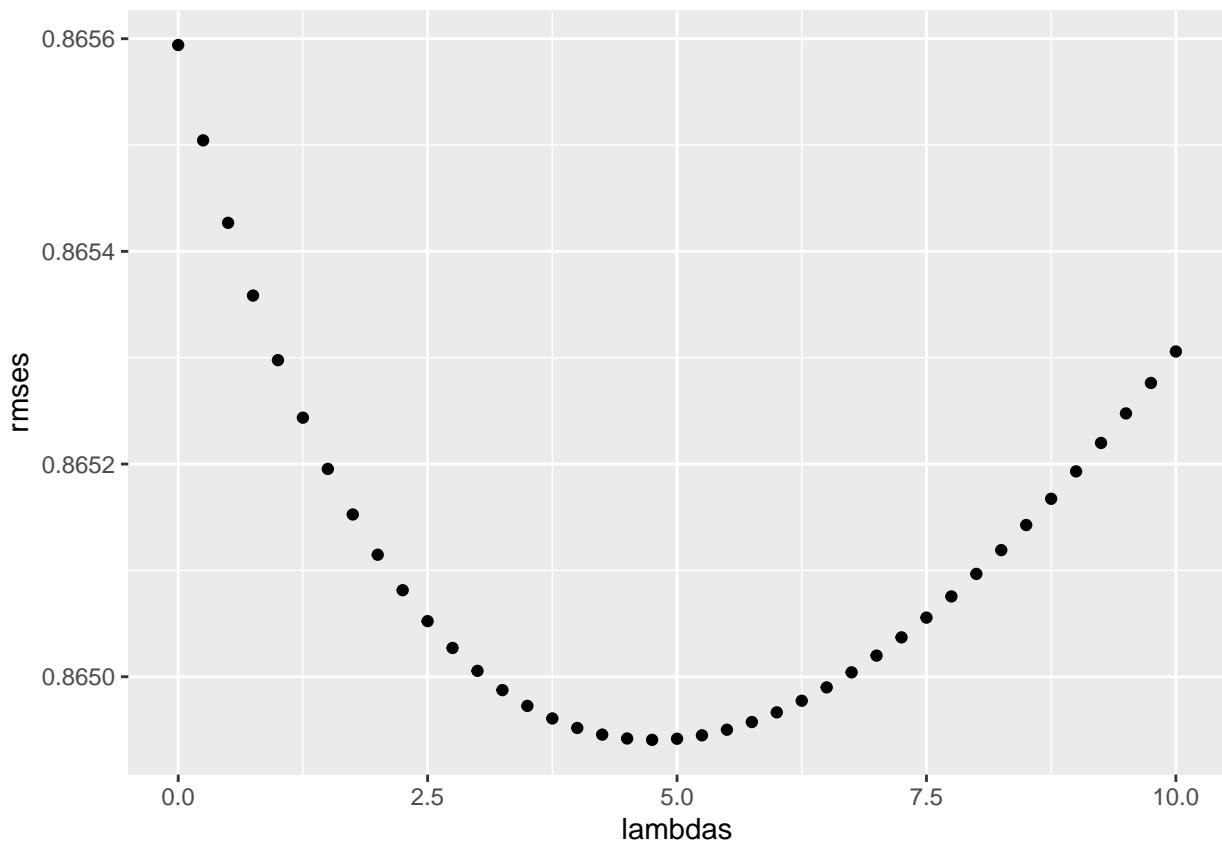
$$\frac{1}{N} \sum_{u,i} (y_{u,i} - \mu - b_i - b_u - b_g)^2 + \lambda (\sum_i b_i^2 + \sum_u b_u^2 + \sum_g b_g^2)$$

The estimates that minimize this can be found similarly to what we did above.

Here **we use 1-fold cross-validation** to pick the optimal λ .

```
lambdas <- seq(0, 10, 0.25)
rmses <- sapply(lambdas, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1))
  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1))
  b_g <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1))
  predicted_ratings <- test_set %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    mutate(pred = mu + b_i + b_u + b_g) %>%
    .$pred
  return(RMSE(predicted_ratings, test_set$rating))
})

qplot(lambdas, rmses)
```



For the Movie + User + Genres model so far, the optimal λ is

```
(lambda <- lambdas[which.min(rmses)])  
## [1] 4.75
```

3.10.1 Results Table Model 1-8

Let's add Regularization to our Model 1-4 results table to get Table 24 and arrange in descending order of RMSE's

Table 24: RMSE Results Models 1-8

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
7	Regularized Movie Effect Model - 1 fold CV	0.9436745
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
8	Regularized Movie + User + Genre Effects Model - 1 fold CV	0.8649406
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330

3.11 Model 9 : Model 6 + Regularization: Choosing the penalty terms

Let's add the penalty term to the model 6 *using 1-fold CV*

```

lambdas <- seq(0, 10, 0.25)
rmses <- sapply(lambdas, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
  b_g <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
  b_d <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    group_by(date) %>%
    summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
  b_r <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    group_by(movieId) %>%
    summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

  predicted_ratings <- test_set %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    left_join(b_r, by='movieId') %>%
    mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
    .$pred

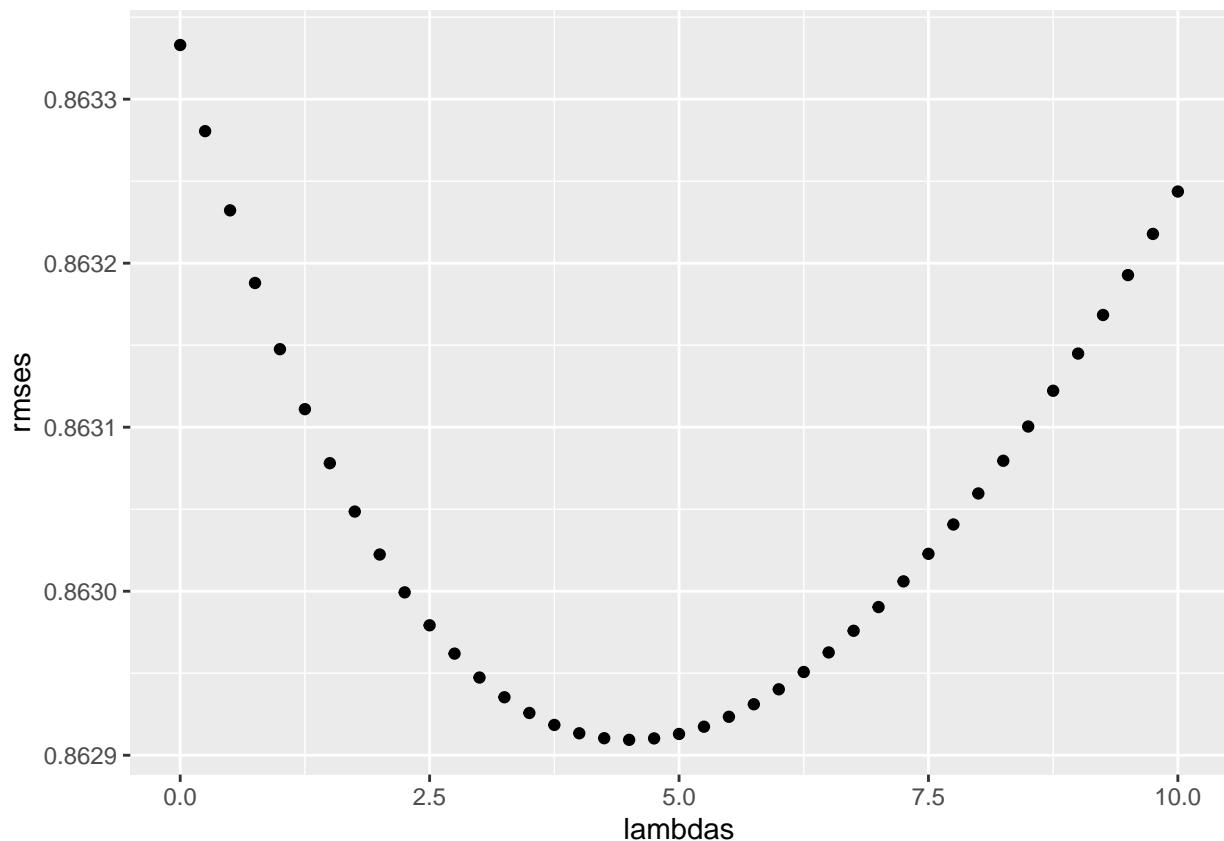
  return(RMSE(predicted_ratings, test_set$rating))
})

(lambda <- lambdas[which.min(rmses)])

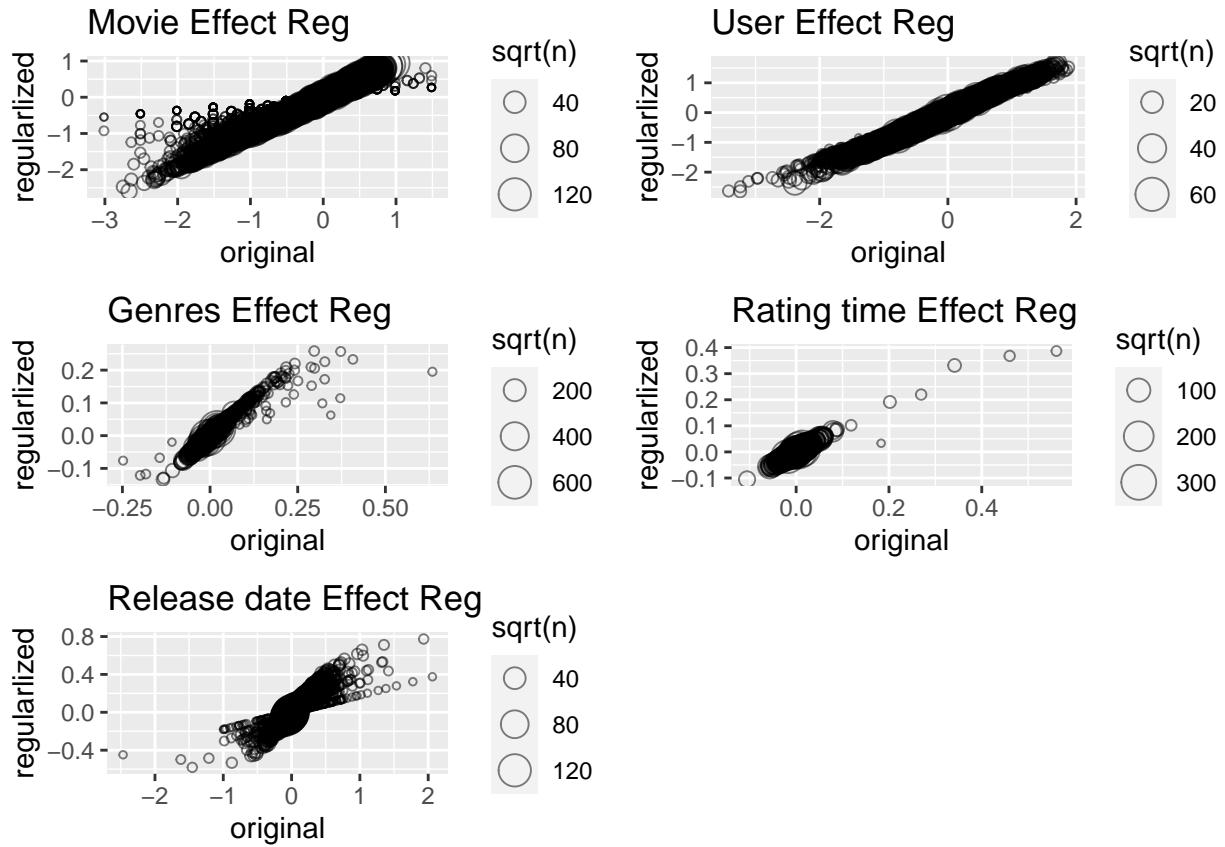
```

[1] 4.5

```
qplot(lambdas, rmses)
```



To see how the estimates shrink, let's make a plot of the regularized estimates versus the least squares estimates. *Here we use only 1 fold*



3.11.1 Results Table Model 1-9

Let's add Regularization to our Model 1-6 results table to get Table 25 and arrange in descending order of RMSE's

Table 25: RMSE Results Models 1-9

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
7	Regularized Movie Effect Model - 1 fold CV	0.9436745
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
8	Regularized Movie + User + Genre Effects Model - 1 fold CV	0.8649406
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330
9	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 1 fold CV	0.8629094

3.12 Multi-fold cross-validation on only edx train_set data set, for the full model 9.

Finally using Multi-fold cross-validation on only our edx train_set data set, for the full model 9, we get the optimal tuning parameter λ , which we will use following this r code chunk against our edx test_set to get our minimum RMSE for the edx data:

```
Sys.time()

[1] "2021-07-28 09:49:07 EDT"

# load("rdas/train_set.rda")

input <- train_set

# k <- 4 # temporary, to be deleted 15-30 min
# k <- 5 # temporary, to be deleted 1 hr, also default for fold()
k <- 10 # temporary, to be deleted 3 hrs
# k <- 20 # temporary, to be deleted 5.45-6 hrs
# k <- 60 #                         19hours
Sys.time()

[1] "2021-07-28 09:49:07 EDT"

set.seed(1, sample.kind="Rounding") # For reproducibility

training_set <- fold(input, k = k, num_col = "rating")
# training_set <- fold(input, k = k, num_col = "rating", num_fold_cols = 2)
# training_set <- fold(input, k = k, cat_col = "genres", id_col = "userId", num_col = "rating")
# rm(input, train_set)
rm(input)

# Order by .folds
training_set <- training_set %>% arrange(.folds)

# training_set %>% kable()
str(training_set)

grouped_df [7,200,043 x 14] (S3: grouped_df/tbl_df/tbl/data.frame)
$ userId      : int [1:7200043] 2 2 2 3 3 3 3 4 4 4 ...
$ movieId     : num [1:7200043] 648 733 1210 1564 1597 ...
$ rating       : num [1:7200043] 2 3 4 4.5 4.5 4 3 5 3 5 ...
$ title        : chr [1:7200043] "Mission: Impossible" "Rock, The" "Star Wars: Episode VI - Return ...
$ genres       : chr [1:7200043] "Action|Adventure|Mystery|Thriller" "Action|Adventure|Thriller" "Act ...
$ rating_date  : POSIXct[1:7200043], format: "1997-07-07 03:04:59" "1997-07-07 03:02:42" ...
$ movie_dt    : num [1:7200043] 1996 1996 1983 1997 1997 ...
$ date         : POSIXct[1:7200043], format: "1997-07-06" "1997-07-06" ...
$ avg_rating   : num [1:7200043] 3.62 3.62 3.62 3.53 3.46 ...
$ avg_rating_rel: num [1:7200043] 3.39 3.7 4 3.57 3.32 ...
$ n           : int [1:7200043] 18992 16175 22584 154 5196 1188 4740 14550 21361 23367 ...
$ years       : num [1:7200043] 22 22 35 21 21 29 16 24 24 28 ...
$ n_year      : num [1:7200043] 863 735 645 7 247 41 296 606 890 835 ...
```

```

$ .folds      : Factor w/ 10 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 ...
- attr(*, "groups")= tibble [10 x 2] (S3: tbl_df/tbl/data.frame)
..$.folds: Factor w/ 10 levels "1","2","3","4",...: 1 2 3 4 5 6 7 8 9 10
..$.rows : list<int> [1:10]
... $. : int [1:720004] 1 2 3 4 5 6 7 8 9 10 ...
... $. : int [1:720004] 720005 720006 720007 720008 720009 720010 720011 720012 720013 720014 ...
... $. : int [1:720005] 1440009 1440010 1440011 1440012 1440013 1440014 1440015 1440016 1440017 1440018 ...
... $. : int [1:720004] 2160014 2160015 2160016 2160017 2160018 2160019 2160020 2160021 2160022 2160023 ...
... $. : int [1:720004] 2880018 2880019 2880020 2880021 2880022 2880023 2880024 2880025 2880026 2880027 ...
... $. : int [1:720005] 3600022 3600023 3600024 3600025 3600026 3600027 3600028 3600029 3600030 3600031 ...
... $. : int [1:720005] 4320027 4320028 4320029 4320030 4320031 4320032 4320033 4320034 4320035 4320036 ...
... $. : int [1:720004] 5040032 5040033 5040034 5040035 5040036 5040037 5040038 5040039 5040040 5040041 ...
... $. : int [1:720004] 5760036 5760037 5760038 5760039 5760040 5760041 5760042 5760043 5760044 5760045 ...
... $. : int [1:720004] 6480040 6480041 6480042 6480043 6480044 6480045 6480046 6480047 6480048 6480049 ...
... @ ptype: int(0)
.. - attr(*, ".drop")= logi TRUE

# lambdas <- seq(0, 10, 0.25)
# we have learnt that the lambda range is 4-7 at this point
lambdas <- seq(4, 7, 0.05)

rmses <- matrix(ncol = length(1:k), nrow = length(lambdas))
rmses <- sapply(1:k, function(fold){

  # Create training set for this iteration
  # Subset all the datapoints where .folds does not match the current fold
  train_set_ts <- training_set[training_set$.folds != fold,]

  # Create test set for this iteration
  # Subset all the datapoints where .folds matches the current fold
  test_set_ts <- training_set[training_set$.folds == fold,]
  test_set_ts <- test_set_ts %>%
    semi_join(train_set_ts, by = "movieId") %>%
    semi_join(train_set_ts, by = "userId") %>%
    semi_join(train_set_ts, by = "date")

  rmses[,fold] <- sapply(lambdas, function(l){
    mu <- mean(train_set_ts$rating)
    b_i <- train_set_ts %>%
      group_by(movieId) %>%
      summarize(b_i = sum(rating - mu)/(n()+1))
    b_u <- train_set_ts %>%
      left_join(b_i, by="movieId") %>%
      group_by(userId) %>%
      summarize(b_u = sum(rating - b_i - mu)/(n()+1))
    b_g <- train_set_ts %>%
      left_join(b_i, by='movieId') %>%
      left_join(b_u, by='userId') %>%
      group_by(genres) %>%
      summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1))
    b_d <- train_set_ts %>%
      left_join(b_i, by='movieId') %>%
      left_join(b_u, by='userId') %>%
      left_join(b_g, by = "genres") %>%
  
```

```

group_by(date) %>%
  summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1))
b_r <- train_set_ts %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1))

predicted_ratings <- test_set_ts %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  left_join(b_r, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .\$pred

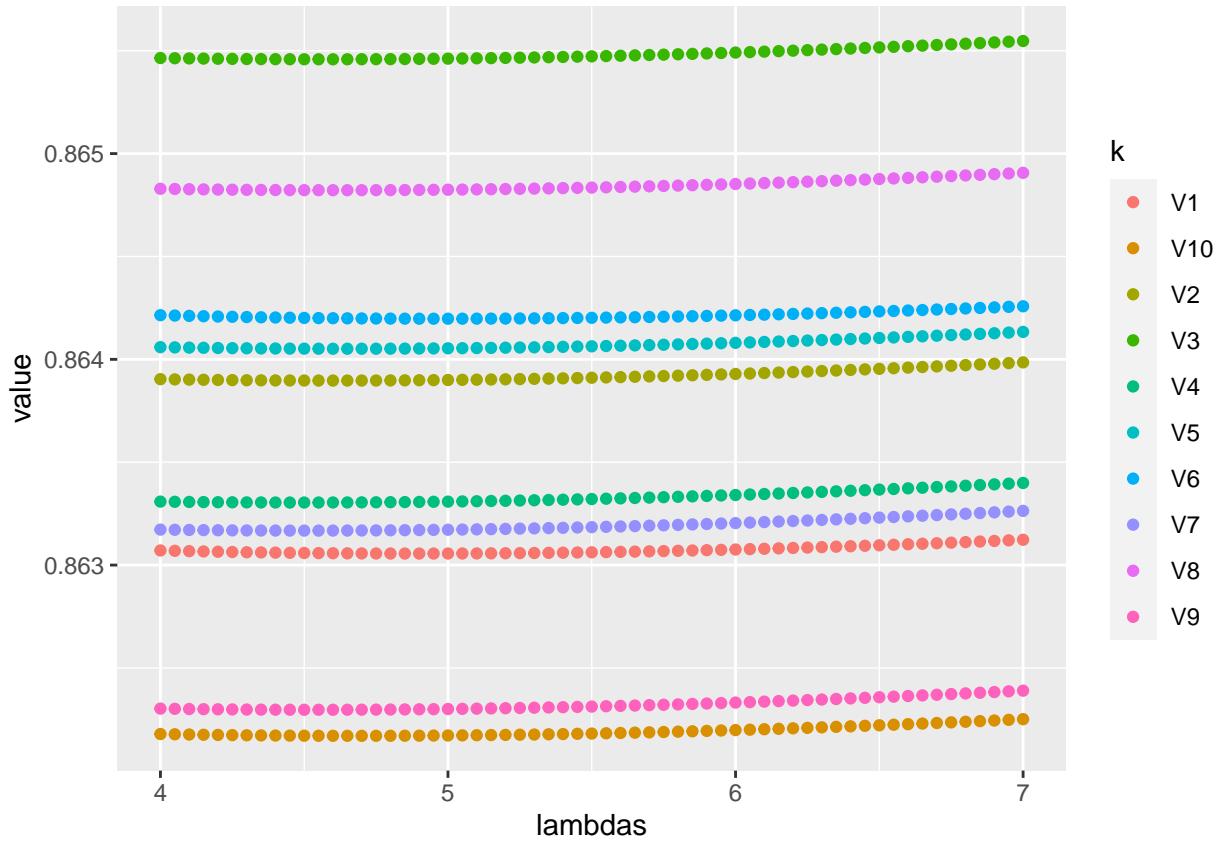
return(RMSE(predicted_ratings, test_set_ts$rating))
})
}
Sys.time()

[1] "2021-07-28 11:46:29 EDT"

rm(training_set,train_set_ts,test_set_ts)
save(rmses,file = "rdas/rmses_k.rda")

rmses_all <- as.data.frame(rmses) %>%
  mutate(lambda=lambdas) %>%
  gather(key = "k", value = "value",1:sum(unique(k)))
rmses_all %>%
  ggplot(aes(lambda,value, col=k)) +
  geom_point()

```



```
# target rmse 0.86490
summary(rmses)
```

V1	V2	V3	V4
Min. :0.8631	Min. :0.8639	Min. :0.8655	Min. :0.8633
1st Qu.:0.8631	1st Qu.:0.8639	1st Qu.:0.8655	1st Qu.:0.8633
Median :0.8631	Median :0.8639	Median :0.8655	Median :0.8633
Mean :0.8631	Mean :0.8639	Mean :0.8655	Mean :0.8633
3rd Qu.:0.8631	3rd Qu.:0.8639	3rd Qu.:0.8655	3rd Qu.:0.8634
Max. :0.8631	Max. :0.8640	Max. :0.8655	Max. :0.8634
V5	V6	V7	V8
Min. :0.8641	Min. :0.8642	Min. :0.8632	Min. :0.8648
1st Qu.:0.8641	1st Qu.:0.8642	1st Qu.:0.8632	1st Qu.:0.8648
Median :0.8641	Median :0.8642	Median :0.8632	Median :0.8648
Mean :0.8641	Mean :0.8642	Mean :0.8632	Mean :0.8648
3rd Qu.:0.8641	3rd Qu.:0.8642	3rd Qu.:0.8632	3rd Qu.:0.8649
Max. :0.8641	Max. :0.8643	Max. :0.8633	Max. :0.8649
V9	V10		
Min. :0.8623	Min. :0.8622		
1st Qu.:0.8623	1st Qu.:0.8622		
Median :0.8623	Median :0.8622		
Mean :0.8623	Mean :0.8622		
3rd Qu.:0.8623	3rd Qu.:0.8622		
Max. :0.8624	Max. :0.8623		

```
summary(rmses_all)
```

```

      lambdas          k           value
Min.    :4.00  Length:610       Min.   :0.8622
1st Qu.:4.75 Class  :character 1st Qu.:0.8631
Median  :5.50 Mode   :character Median  :0.8636
Mean    :5.50                   Mean   :0.8637
3rd Qu.:6.25                   3rd Qu.:0.8642
Max.    :7.00                   Max.   :0.8655

# Get lambda for min rmse values for each fold
lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]

[1] 4.90 4.60 4.60 4.50 4.65 5.00 4.50 4.60 4.55 4.70

rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]

[1] 4.90 4.60 4.60 4.50 4.65 5.00 4.50 4.60 4.55 4.70

# Get mean lambda from min rmse values for each fold
mean(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

[1] 4.66

mean(rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

[1] 4.66

(lambda_opt <- mean(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]))

[1] 4.66

# Get median lambda from min rmse values for each fold
median(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

[1] 4.6

median(rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

[1] 4.6

# Get range/(min,max) lambda from min rmse values for each fold
range(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

[1] 4.5 5.0

save(rmses_all,file = "rdas/rmses_all_k.rda")

Sys.time()

[1] "2021-07-28 11:46:30 EDT"

```

3.13 Model 10 : Model 6 + Regularization from multi-fold CV(Model 9 with multi-fold CV)

We will use the penalty term from our previous section where we got our optimal lambda_opt 4.66 using 10-fold CV to get minimum RMSE for the edx test_set

```
# load("rdas/train_set.rda")
# load("rdas/test_set.rda")

rmses_min <- sapply(lambda_opt, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
  b_g <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
  b_d <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    group_by(date) %>%
    summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
  b_r <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    group_by(movieId) %>%
    summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

  predicted_ratings <- test_set %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    left_join(b_r, by='movieId') %>%
    mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
    .$pred

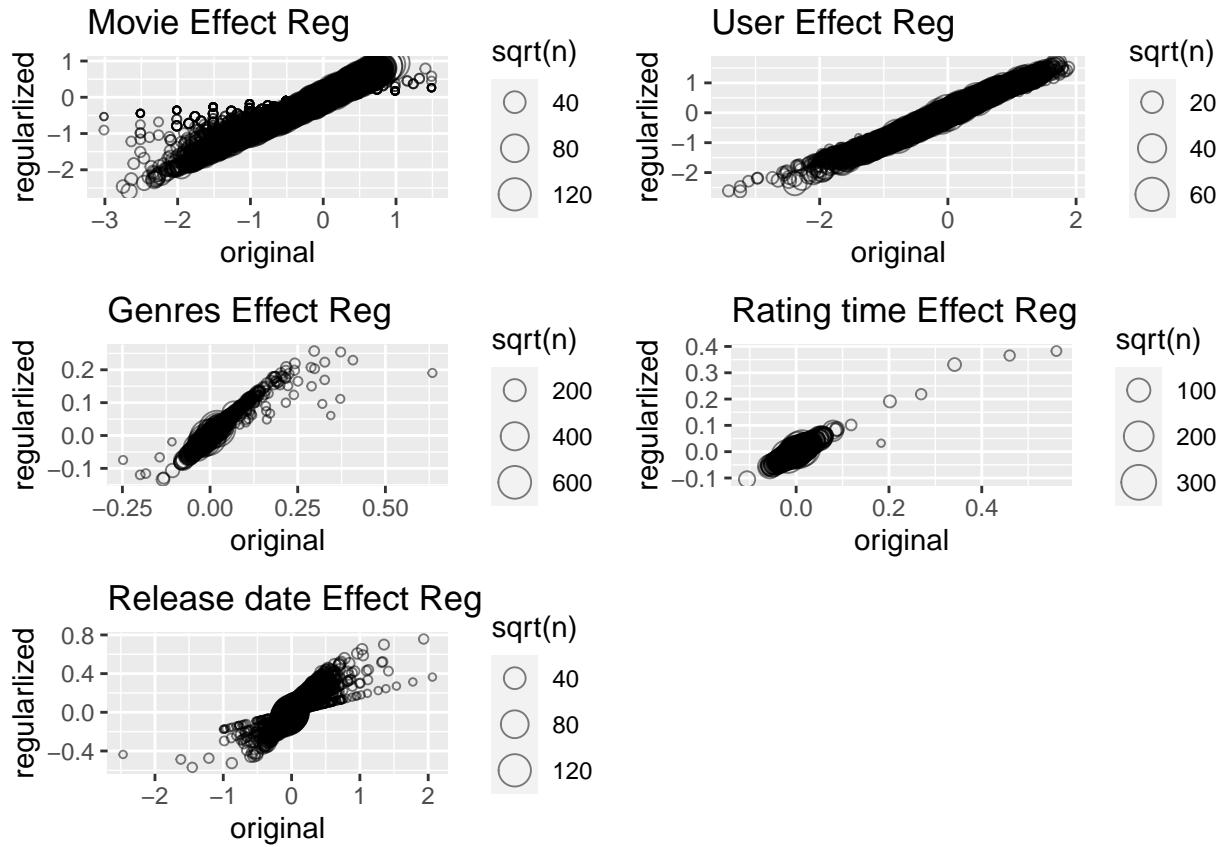
  return(RMSE(predicted_ratings, test_set$rating))
})

(rmses_min)

[1] 0.8629097

save(rmses_min, file = "rdas/rmses_min.rda")
```

To see how the estimates shrink, let's make a plot of the regularized estimates versus the least squares estimates. *Here we use 10 fold*



3.13.1 Results Table Model 1-10

Let's add Model 10 results table to get Table 26 and arrange in descending order of RMSE's

Table 26: RMSE Results Models 1-10

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
7	Regularized Movie Effect Model - 1 fold CV	0.9436745
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
8	Regularized Movie + User + Genre Effects Model - 1 fold CV	0.8649406
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330
10	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 10 fold CV	0.8629097
9	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 1 fold CV	0.8629094

3.14 Final Hold-Out test of my final algorithm

I have developed my algorithm using the edx set. For a final test of my final algorithm, I predict movie ratings in the validation set (the final hold-out test set) as if they were unknown. RMSE⁸ will be used to evaluate how close my predictions are to the true values in the validation set (the final hold-out test set). My target is RMSE < 0.86490.

```
load("rdas/edx_md.rda")
load("rdas/validation_md.rda")
(lambda_opt)

[1] 4.66

rmses_val <- sapply(lambda_opt, function(l){
  mu <- mean(edx_md$rating)
  b_i <- edx_md %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
  b_u <- edx_md %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
  b_g <- edx_md %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
  b_d <- edx_md %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    group_by(date) %>%
    summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
  b_r <- edx_md %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    group_by(movieId) %>%
    summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

  predicted_ratings <- validation_md %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    left_join(b_r, by='movieId') %>%
    mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
    .$pred

  return(RMSE(predicted_ratings, validation_md$rating))
})
```

⁸https://en.wikipedia.org/wiki/Root-mean-square_deviation

```
(rmses_val)

[1] 0.8634961

save(rmses_val, file = "rdas/rmses_val.rda")
```

3.14.1 Results Table Model 1-10

Let's add Model 10 results table to get Table 27 and arrange in descending order of RMSE's

Table 27: RMSE Results Models 1-10 Validation

Index	Method	RMSE
1	Just the average	1.0599043
2	Movie Effect Model	0.9437429
7	Regularized Movie Effect Model - 1 fold CV	0.9436745
3	Movie + User Effects Model	0.8659320
4	Movie + User + Genres Effects Model	0.8655941
5	Movie + User + Genres + Rating Time Effects Model	0.8654205
8	Regularized Movie + User + Genre Effects Model - 1 fold CV	0.8649406
VAL	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 10 fold CV - Validation	0.8634961
6	Movie + User + Genres + Rating Time + Release date Effects Model	0.8633330
10	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 10 fold CV	0.8629097
9	Regularized Movie + User + Genre + Rating Time + Release date Effect Model - 1 fold CV	0.8629094

4 Conclusion

In section 1.1 I set out with a goal to develop my algorithm using the edx set (MovieLens data). For a final test of my final algorithm, I had to predict movie ratings in the validation set (the final hold-out test set) as if they were unknown. RMSE was to be used to evaluate how close my predictions were to the true values in the validation set (the final hold-out test set). My target was $\text{RMSE} < 0.86490$.

In Section 3.14, using models - “Regularized Movie + User + Genre + Rating Time + Release date Effect Model” I achieved my goal with an RMSE of 0.8634961.

The lesson here is that having lots of models is useful for the incremental results needed to win competitions, but practically, excellent systems can be built with just a few well-selected models.

For the future, I will attempt to significantly simplify the process of selecting models and methods/techniques. For a given dataset, my goal will be to find algorithms that achieve close to maximum accuracy while minimizing computation time required for training.

5 Appendix: All code for this report

```
knitr::knit_hooks$set(time_it = local({
  now <- NULL
  function(before, options) {
    if (before) {
      # record the current time before each chunk
      now <<- Sys.time()
    } else {
      # calculate the time difference after a chunk
      res <- difftime(Sys.time(), now)
      # return a character string to show the time
      # paste("Time for this code chunk to run:", res)
      paste("Time for the chunk", options$label, "to run:", res)
    }
  }
}))
```

```
# knit_hooks$get("inline")
# knitr::opts_chunk$set(fig.pos = "!H", out.extra = "")
knitr::opts_chunk$set(echo = TRUE,
                     fig.path = "figures/")

# Beware, using the "time_it" hook messes up fig.cap, \label, \ref
# knitr::opts_chunk$set(time_it = TRUE)
#knitr::opts_chunk$set(eval = FALSE)
library(ggplot2)
library(kableExtra)

#####
# Create edx set, validation set
#####

# Note: this process could take a couple of minutes

if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
if(!require(lubridate)) install.packages("lubridate", repos = "http://cran.us.r-project.org")
if(!require(groupdata2)) install.packages("groupdata2", repos = "http://cran.us.r-project.org")
if(!require(gridExtra)) install.packages("gridExtra", repos = "http://cran.us.r-project.org")

# https://www.tidyverse.org/blog/2020/05/dplyr-1-0-0-last-minute-additions/
options(dplyr.summarise.inform = FALSE)

# MovieLens 10M dataset:
# https://grouplens.org/datasets/movielens/10m/
# http://files.grouplens.org/datasets/movielens/ml-10m-README.html
# http://files.grouplens.org/datasets/movielens/ml-10m.zip

dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
```

```

ratings <- fread(text = gsub("::", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
                  col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\::", 3)
colnames(movies) <- c("movieId", "title", "genres")
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
                                              title = as.character(title),
                                              genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind="Rounding")
# if using R 3.5 or earlier, use `set.seed(1)` instead
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

save(ratings, file = "rdas/ratings.rda")
save(movies, file = "rdas/movies.rda")
rm(dl, ratings, movies, test_index, temp, movielens, removed)

save(edx, file = "rdas/edx.rda")
save(validation, file = "rdas/validation.rda")
kable(edx[1:10], "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Movielens data\\label{tb1}")
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
unique_u_m_g <- edx %>%
  summarize(unique_users = n_distinct(userId),
            unique_movies = n_distinct(movieId),
            unique_genres = n_distinct(genres))

kable(unique_u_m_g, "latex",
      booktabs=TRUE, linesep="",
      caption="Unique Users, Movies and Genres\\label{tbl:uniq_users_movies_genres}") %>% kable_styling()
keep <- edx %>%
  dplyr::count(movieId) %>%
  top_n(4) %>%
  pull(movieId)
tab <- edx %>%
  filter(userId %in% c(13:20)) %>%
  filter(movieId %in% keep) %>%
  select(userId, title, rating) %>%
  spread(title, rating)

kable(tab, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Matrix of seven users and four movies\\label{tbl:matrix}")

```

```

    kable_styling(latex_options=c("HOLD_position"), font_size=8)
users <- sample(unique(edx$userId), 100)
rafaelib::mypar()
edx %>% filter(userId %in% users) %>%
  select(userId, movieId, rating) %>%
  mutate(rating = 1) %>%
  spread(movieId, rating) %>% select(sample(ncol(.), 100)) %>%
  as.matrix() %>% t(.) %>%
  image(1:100, 1:100, , xlab="Movies", ylab="Users")
abline(h=0:100+0.5, v=0:100+0.5, col = "grey")
edx %>%
  dplyr::count(movieId) %>%
  ggplot(aes(n)) +
  geom_histogram(bins = 30, color = "black") +
  scale_x_log10() + # try with and without this line
  ggtitle("Movies")
edx %>%
  dplyr::count(userId) %>%
  ggplot(aes(n)) +
  geom_histogram(bins = 30, color = "black") +
  scale_x_log10() + # try with and without this line
  ggtitle("Users")

rm(tab,keep,unique_u_m_g,users)
edx_dt <- edx %>%
  mutate(rating_date = as_datetime(timestamp)) %>%
  select(-timestamp)
rm(edx)

kable(edx_dt[1:5,], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens edx data with ratings")
  kable_styling(latex_options=c("HOLD_position"), font_size=8)
td <- edx_dt$title
movie_dt <- str_extract(td, "\\([0-9]{4}\\)$") %>%
  str_extract("[0-9]{4}")
title_o <- str_replace(td, "\\([0-9]{4}\\)$", "")
edx_md <- edx_dt %>% mutate(title=title_o, movie_dt=as.numeric(movie_dt))
rm(td,movie_dt,title_o,edx_dt)
save(edx_md, file = "rdas/edx_md.rda")

kable(edx_md[1:5,], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens edx data with movie IDs")
  kable_styling(latex_options=c("HOLD_position"), font_size=8)
validation_dt <- validation %>%
  mutate(rating_date = as_datetime(timestamp)) %>%
  select(-timestamp)
rm(validation)

kable(validation_dt[1:5,], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens validation data")
  kable_styling(latex_options=c("HOLD_position"), font_size=8)
td <- validation_dt$title
movie_dt <- str_extract(td, "\\([0-9]{4}\\)$") %>%
  str_extract("[0-9]{4}")
title_o <- str_replace(td, "\\([0-9]{4}\\)$", "")
validation_md <- validation_dt %>% mutate(title=title_o, movie_dt=as.numeric(movie_dt))
rm(td,movie_dt,title_o,validation_dt)

```

```

save(validation_md, file = "rdas/validation_md.rda")

kable(validation_md[1:5], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens validation data with avg rating")
  kable_styling(latex_options=c("HOLD_position"), font_size=8)

# rm(validation_md)
edx_md %>% group_by(genres) %>%
  summarize(n = n(), avg = mean(rating), se = sd(rating)/sqrt(n())) %>%
  filter(n >= 1000) %>%
  mutate(genres = reorder(genres, avg)) %>%
  ggplot(aes(x = genres, y = avg, ymin = avg - 2*se, ymax = avg + 2*se)) +
  geom_point() +
  geom_errorbar() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
edx_md %>% mutate(date = round_date(rating_date, unit = "week")) %>%
  group_by(date) %>%
  summarize(avg_rating = mean(rating)) %>%
  ggplot(aes(date, avg_rating)) +
  geom_point() +
  geom_smooth()
edx_md <- edx_md %>% mutate(date = round_date(rating_date, unit = "week")) %>%
  group_by(date) %>%
  mutate(avg_rating = mean(rating)) %>% ungroup()
save(edx_md, file = "rdas/edx_md.rda")

kable(edx_md[1:5], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens edx data with avg rating")
  kable_styling(latex_options=c("HOLD_position"), font_size=6)
validation_md <- validation_md %>% mutate(date = round_date(rating_date, unit = "week")) %>%
  group_by(date) %>%
  mutate(avg_rating = mean(rating)) %>% ungroup()
save(validation_md, file = "rdas/validation_md.rda")

kable(validation_md[1:5], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens validation data with avg rating")
  kable_styling(latex_options=c("HOLD_position"), font_size=6)

# rm(validation_md)
ratings_m_y <- edx_md %>%
  group_by(movieId) %>%
  summarize(n = n_distinct(userId), year = as.character(first(movie_dt))) %>%
  mutate(sqrt_n=sqrt(n)) %>%
  group_by(year) %>%
  mutate(r_median=median(sqrt_n)) %>% select(year,r_median) %>%
  unique() %>% arrange(desc(r_median))

# min(edx_md$movie_dt)
# # [1] 1915
# max(edx_md$movie_dt)
# # [1] 2008

# View(ratings_m_y)
ggplot(data=ratings_m_y, aes(x=year, y=r_median) ) +
  geom_point( size=1, colour="#000080" ) +
  theme_light(base_size=8) +
  # labs(title="", x="\nyear", y="r_median\n") +

```

```

scale_x_discrete(breaks=seq(min(edx_md$movie_dt),max(edx_md$movie_dt),by=2)) +
theme(axis.text.x = element_text(angle = 90, hjust = 1))

spline_int <- as.data.frame(spline(ratings_m_y$year, ratings_m_y$r_median))

ggplot(data=ratings_m_y, aes(x=year, y=r_median) ) +
geom_point(data = spline_int, aes(x = x, y = y)) +
geom_line(data = spline_int, aes(x = x, y = y)) +
theme_light(base_size=8) +
# scale_x_discrete(breaks=seq(min(edx_md$movie_dt),max(edx_md$movie_dt),by=2)) +
theme(axis.text.x = element_text(angle = 90, hjust = 1))

rm(ratings_m_y)

res2_1993_plus <- edx_md %>%
filter(movie_dt >= 1993) %>%
group_by(movieId) %>%
# summarize(avg_rating = mean(rating), n = n(), title=title[1], years=2018 - min(movie_dt)) %>%
summarize(avg_rating = mean(rating), n = n(), title=title[1], years=2018 - first(movie_dt)) %>%
mutate(n_year = n / years) %>%
top_n(25, n_year) %>%
arrange(desc(n_year))
# qplot(res2_1993_plus$years,res2_1993_plus$n)
ggplot(res2_1993_plus,aes(years,n)) +
geom_line() +
geom_smooth() +
geom_vline(xintercept = 25) # 1993

rm(res2_1993_plus)
res3 <- edx_md %>%
filter(movie_dt >= 1993) %>%
group_by(movieId) %>%
summarize(avg_rating = mean(rating), n = n(), title=title[1], years=2018 - min(movie_dt)) %>%
mutate(n_year = n / years)
res3 %>%
group_by(round(n_year)) %>%
ggplot(aes(n_year,avg_rating)) + geom_point() + geom_smooth()

rm(res3)
res3_b <- edx_md %>%
filter(movie_dt < 1993) %>%
group_by(movieId) %>%
summarize(avg_rating = mean(rating), n = n(), title=title[1], years=2018 - min(movie_dt)) %>%
mutate(n_year = n / years)
res3_b %>%
group_by(round(n_year)) %>%
ggplot(aes(n_year,avg_rating)) + geom_point() + geom_smooth()

rm(res3_b)
res3_all_edx <- edx_md %>%
group_by(movieId) %>%
summarize(avg_rating_rel = mean(rating), n = n(), title=title[1], years=2018 - first(movie_dt)) %>%
mutate(n_year = round(n / years))
res3_all_edx %>%

```

```

group_by(n_year) %>%
  ggplot(aes(n_year, avg_rating_rel)) + geom_point() + geom_smooth()
save(res3_all_edx, file = "rdas/res3_all_edx.rda")
res3_all_edx_2 <- res3_all_edx %>% select(-title)
edx_md <- edx_md %>% left_join(res3_all_edx_2, by='movieId')
save(edx_md, file = "rdas/edx_md.rda")

kable(edx_md[1:5,], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens edx data with a
  kable_styling(latex_options=c("HOLD_position"), font_size=4)
str(edx_md)
rm(res3_all_edx,res3_all_edx_2)
res3_all_val <- validation_md %>%
  group_by(movieId) %>%
  summarize(avg_rating_rel = mean(rating), n = n(), title=title[1], years=2018 - first(movie_dt)) %>%
  mutate(n_year = round(n / years))
# res3_all_val %>%
#   group_by(n_year) %>%
#   ggplot(aes(n_year, avg_rating_rel)) + geom_point() + geom_smooth()
# save(res3_all_val, file = "rdas/res3_all_val.rda")

res3_all_val_2 <- res3_all_val %>% select(-title)
validation_md <- validation_md %>% left_join(res3_all_val_2, by='movieId')
save(validation_md, file = "rdas/validation_md.rda")

kable(validation_md[1:5,], "latex", escape=TRUE, booktabs=TRUE, linesep="", caption="Movielens validation
  kable_styling(latex_options=c("HOLD_position"), font_size=4)
str(validation_md)
rm(res3_all_val,res3_all_val_2,validation_md)
# Sys.time()
set.seed(1, sample.kind="Rounding")
test_index <- createDataPartition(y = edx_md$rating, times = 1,
                                   p = 0.2, list = FALSE)
train_set <- edx_md[-test_index,]
test_set <- edx_md[test_index,]

# To make sure we don't include users and movies in the test set that do not
# appear in the training set, we remove these entries using the semi_join
# function:

test_set <- test_set %>%
  semi_join(train_set, by = "movieId") %>%
  semi_join(train_set, by = "userId")
RMSE <- function(true_ratings, predicted_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}

save(test_index, file = "rdas/test_index.rda")
save(train_set, file = "rdas/train_set.rda")
save(test_set, file = "rdas/test_set.rda")

rm(edx_md, test_index)
(mu_hat <- mean(train_set$rating))
(model_1_rmse <- RMSE(test_set$rating, mu_hat))
predictions <- rep(2.5, nrow(test_set))

```

```

RMSE(test_set$rating, predictions)

predictions <- rep(3, nrow(test_set))
RMSE(test_set$rating, predictions)

predictions <- rep(4, nrow(test_set))
RMSE(test_set$rating, predictions)

rm(predictions)
rmse_results <- tibble(Index = "1", Method = "Just the average", RMSE = model_1_rmse)

save(rmse_results, file = "rdas/rmse_results.rda")
save(model_1_rmse, file = "rdas/model_1_rmse.rda")

rm(mu_hat, model_1_rmse)

# rmse_results %>% knitr::kable()
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Model 1\\\
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
mu <- mean(train_set$rating)

movie_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = mean(rating - mu))

save(mu, file = "rdas/mu.rda")
save(movie_avgs, file = "rdas/movie_avgs.rda")
movie_avgs %>%
  ggplot(aes(b_i)) +
  geom_histogram(bins = 10, color = "black") +
  # scale_x_log10() + # try with and without this line
  ggtitle("Movie effect or bias distribution")
predicted_ratings_model_2 <- mu + test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  .$b_i

(model_2_rmse <- RMSE(predicted_ratings_model_2, test_set$rating))

save(predicted_ratings_model_2, file = "rdas/predicted_ratings_model_2.rda")
save(model_2_rmse, file = "rdas/model_2_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "2", Method="Movie Effect Model",
                                 RMSE = model_2_rmse))

save(rmse_results, file = "rdas/rmse_results.rda")
rm(model_2_rmse, predicted_ratings_model_2)
# rmse_results %>% knitr::kable()
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-2\\\
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
train_set %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating)) %>%
  filter(n()>=100) %>%

```

```

ggplot(aes(b_u)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for users who have rated over 100 movies")
train_set %>%
  group_by(userId) %>%
  summarise(b_u = mean(rating)) %>%
  ggplot(aes(b_u)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for users who have rated any movies")
user_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  group_by(userId) %>%
  summarise(b_u = mean(rating - mu - b_i))

save(user_avgs, file = "rdas/user_avgs.rda")
user_avgs %>%
  ggplot(aes(b_u)) +
  geom_histogram(bins = 10, color = "black") +
  # scale_x_log10() + # try with and without this line
  ggtitle("User effect or bias distribution")
predicted_ratings_model_3 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>%
  .$pred

(model_3_rmse <- RMSE(predicted_ratings_model_3, test_set$rating))

save(predicted_ratings_model_3, file = "rdas/predicted_ratings_model_3.rda")
save(model_3_rmse, file = "rdas/model_3_rmse.rda")
rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "3", Method="Movie + User Effects Model",
                                  RMSE = model_3_rmse))
save(rmse_results, file = "rdas/rmse_results.rda")
rm(model_3_rmse, predicted_ratings_model_3)
# rmse_results %>% knitr::kable()
  kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-3")
    kable_styling(latex_options=c("HOLD_position"), font_size=7)
train_set %>%
  group_by(genres) %>%
  summarise(mu_g = mean(rating)) %>%
  ggplot(aes(mu_g)) +
  geom_histogram(bins = 30, color = "black") +
  ggtitle("Average rating for movies of category genres")
q8_1 <- train_set %>%
  group_by(movieId) %>%
  summarise(n = n(), genres=genres[1])

q8_2 <- q8_1 %>% group_by(genres) %>%
  summarise(nr=sum(n)) %>%
  filter(nr>1000)

q8 <- train_set %>%
  group_by(genres) %>%

```

```

summarize(n = n(), avg = mean(rating), se = sd(rating)/sqrt(n())) %>%
filter(n >= 1000)

q8 %>%
  mutate(genres = reorder(genres, avg)) %>%
  ggplot(aes(x = genres, y = avg, ymin = avg - 2*se, ymax = avg + 2*se)) +
  geom_point() +
  geom_errorbar() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))

mean(q8$avg)

median(q8$avg)

rm(q8_1,q8_2,q8)
genres_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = mean(rating - mu - b_i - b_u))
genres_avgs %>%
  ggplot(aes(b_g)) +
  geom_histogram(bins = 10, color = "black") +
  # scale_x_log10() + # try with and without this line
  ggtitle("Genres effect or bias distribution")
predicted_ratings_model_4 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  mutate(pred = mu + b_i + b_u + b_g) %>%
  .$pred

(model_4_rmse <- RMSE(predicted_ratings_model_4, test_set$rating))

save(genres_avgs, file = "rdas/genres_avgs.rda")
save(predicted_ratings_model_4, file = "rdas/predicted_ratings_model_4.rda")
save(model_4_rmse, file = "rdas/model_4_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "4", Method="Movie + User + Genres Effects Model",
                                  RMSE = model_4_rmse))

save(rmse_results, file = "rdas/rmse_results.rda")

rm(model_4_rmse, predicted_ratings_model_4)
# rmse_results %>% knitr::kable()
  kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-")
    kable_styling(latex_options=c("HOLD_position"), font_size=7)
time_effect_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  group_by(date) %%%
  summarize(b_d = mean(avg_rating - mu - b_i - b_u - b_g))

```

```

time_effect_avgs %>%
  ggplot(aes(b_d)) +
  geom_histogram(bins = 10, color = "black") +
  # scale_x_log10() + # try with and without this line
  ggtitle("Rating time effect or bias distribution")
predicted_ratings_model_5 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  left_join(time_effect_avgs, by = "date") %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d) %>%
  .$pred

(model_5_rmse <- RMSE(predicted_ratings_model_5, test_set$rating))

save(time_effect_avgs, file = "rdas/time_effect_avgs.rda")
save(predicted_ratings_model_5, file = "rdas/predicted_ratings_model_5.rda")
save(model_5_rmse, file = "rdas/model_5_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "5", Method="Movie + User + Genres + Rating Time Effects Model",
                                  RMSE = model_5_rmse))

save(rmse_results, file = "rdas/rmse_results.rda")

rm(model_5_rmse, predicted_ratings_model_5)
# rmse_results %>% knitr::kable()
  kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-5")
    kable_styling(latex_options=c("HOLD_position"), font_size=7)

rel_effect_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  left_join(time_effect_avgs, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = mean(avg_rating_rel - mu - b_i - b_u - b_g - b_d))

save(rel_effect_avgs, file = "rdas/rel_effect_avgs.rda")
rel_effect_avgs %>%
  ggplot(aes(b_r)) +
  geom_histogram(bins = 10, color = "black") +
  # scale_x_log10() + # try with and without this line
  ggtitle("Release Date effect or bias distribution")
predicted_ratings_model_6 <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by='genres') %>%
  left_join(time_effect_avgs, by = "date") %>%
  left_join(rel_effect_avgs, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .$pred

(model_6_rmse <- RMSE(predicted_ratings_model_6, test_set$rating))

```

```

save(predicted_ratings_model_6, file = "rdas/predicted_ratings_model_6.rda")
save(model_6_rmse, file = "rdas/model_6_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "6", Method="Movie + User + Genres + Rating Time + Release date",
                                  RMSE = model_6_rmse))

rmse_results <- rmse_results %>% arrange(desc(RMSE))

save(rmse_results, file = "rdas/rmse_results.rda")
# rmse_results %>% knitr::kable()
  kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-6")
    kable_styling(latex_options=c("HOLD_position"), font_size=7)
test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  mutate(residual = rating - (mu + b_i)) %>%
  arrange(desc(abs(residual))) %>%
  slice(1:10) %>%
  pull(title) %>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="",
        caption="Without Regularization 10 largest mistakes\\label{tbl:without_regularization_10_largest}")
    kable_styling(latex_options=c("HOLD_position"))
load("rdas/edx.rda")
movie_titles <- edx %>%
  select(movieId, title) %>%
  distinct()
save(movie_titles, file = "rdas/movie_titles.rda")
rm(edx)
movie_avgs %>% left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  slice(1:10) %>%
  pull(title)%>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Without Regularization 10 Best movies")
    kable_styling(latex_options=c("HOLD_position"))
movie_avgs %>% left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %%%
  slice(1:10) %%%
  pull(title)%>%
  kable("latex", escape=FALSE, booktabs=TRUE, linesep="", caption="Without Regularization 10 Worst movies")
    kable_styling(latex_options=c("HOLD_position"))
train_set %>% count(movieId) %>%
  left_join(movie_avgs, by="movieId") %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  slice(1:10) %>%
  pull(n)
train_set %>% count(movieId) %>%
  left_join(movie_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %>%
  slice(1:10) %>%
  pull(n)
lambdas <- seq(0, 10, 0.25)
mu <- mean(train_set$rating)

```

```

just_the_sum <- train_set %>%
  group_by(movieId) %>%
  summarize(s = sum(rating - mu), n_i = n())
rmses <- sapply(lambdas, function(l){
  predicted_ratings <- test_set %>%
    left_join(just_the_sum, by='movieId') %>%
    mutate(b_i = s/(n_i+l)) %>%
    mutate(pred = mu + b_i) %>%
    .$pred
  return(RMSE(predicted_ratings, test_set$rating))
})

(lambda <- lambdas[which.min(rmses)])
qplot(lambdas, rmses)

mu <- mean(train_set$rating)
movie_reg_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda), n_i = n())
rm(just_the_sum)
tibble(original = movie_avgs$b_i,
       regularized = movie_reg_avgs$b_i,
       n = movie_reg_avgs$n_i) %>%
  ggplot(aes(original, regularized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Movie Effect Reg")
train_set %>%
  dplyr::count(movieId) %>%
  left_join(movie_reg_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(desc(b_i)) %>%
  select(title, b_i, n) %>%
  slice(1:10) %>%
  kable("latex", escape=TRUE, booktabs=TRUE, linesep="", caption="With Regularization top 10 best movies")
train_set %>%
  dplyr::count(movieId) %>%
  left_join(movie_reg_avgs) %>%
  left_join(movie_titles, by="movieId") %>%
  arrange(b_i) %>%
  select(title, b_i, n) %>%
  slice(1:10) %>%
  kable("latex", escape=TRUE, booktabs=TRUE, linesep="", caption="With Regularization top 10 worst movies")
  kable_styling(latex_options=c("HOLD_position"))
predicted_ratings_model_7 <- test_set %>%
  left_join(movie_reg_avgs, by='movieId') %>%
  mutate(pred = mu + b_i) %>%
  .$pred

(model_7_rmse <- RMSE(predicted_ratings_model_7, test_set$rating))
save(model_7_rmse, file = "rdas/model_7_rmse.rda")
#
# ````
# ````{r, e_u pls_4, echo=FALSE, eval=TRUE}

```

```

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "7", Method="Regularized Movie Effect Model - 1 fold CV",
                                  RMSE = model_7_rmse))

save(rmse_results, file = "rdas/rmse_results.rda")

rm(model_7_rmse, predicted_ratings_model_7, movie_reg_avgs, lambda)
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-7")
kable_styling(latex_options=c("HOLD_position"), font_size=7)
lambdas <- seq(0, 10, 0.25)
rmses <- sapply(lambdas, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1))
  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1))
  b_g <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1))
  predicted_ratings <- test_set %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    mutate(pred = mu + b_i + b_u + b_g) %>%
    .$pred
  return(RMSE(predicted_ratings, test_set$rating))
})

qplot(lambdas, rmses)
(lambda <- lambdas[which.min(rmses)])
save(rmses, file = "rdas/rmses8.rda")
model_8_rmse <- min(rmses)
save(model_8_rmse, file = "rdas/model_8_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "8", Method="Regularized Movie + User + Genre Effects Model - 8",
                                  RMSE = model_8_rmse))

rmse_results <- rmse_results %>% arrange(desc(RMSE))

save(rmse_results, file = "rdas/rmse_results.rda")

rm(lambda, lambdas, rmses)
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-8")
kable_styling(latex_options=c("HOLD_position"), font_size=7)
lambdas <- seq(0, 10, 0.25)
rmses <- sapply(lambdas, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%

```

```

group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
b_u <- train_set %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
b_g <- train_set %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
b_d <- train_set %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  group_by(date) %>%
  summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
b_r <- train_set %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

predicted_ratings <- test_set %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  left_join(b_r, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .$pred

return(RMSE(predicted_ratings, test_set$rating))
}

(lambda <- lambdas[which.min(rmses)])
qplot(lambdas, rmses)

mu <- mean(train_set$rating)
movie_reg_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda), n_i = n())
user_reg_avgs <- train_set %>%
  left_join(movie_avgs, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda), n_i = n())
genres_reg_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+lambda), n_i = n())
time_effect_reg_avgs <- train_set %>%

```

```

left_join(movie_avgs, by='movieId') %>%
left_join(user_avgs, by='userId') %>%
left_join(genres_avgs, by = "genres") %>%
  group_by(date) %>%
  summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n() + lambda), n_i = n())
rel_effect_reg_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  left_join(time_effect_avgs, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n() + lambda), n_i = n())

mp <- tibble(original = movie_avgs$b_i,
             regularlized = movie_reg_avgs$b_i,
             n = movie_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Movie Effect Reg")
up <- tibble(original = user_avgs$b_u,
             regularlized = user_reg_avgs$b_u,
             n = user_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("User Effect Reg")
gp <- tibble(original = genres_avgs$b_g,
             regularlized = genres_reg_avgs$b_g,
             n = genres_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Genres Effect Reg")
tp <- tibble(original = time_effect_avgs$b_d,
             regularlized = time_effect_reg_avgs$b_d,
             n = time_effect_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Rating time Effect Reg")
rp <- tibble(original = rel_effect_avgs$b_r,
             regularlized = rel_effect_reg_avgs$b_r,
             n = rel_effect_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Release date Effect Reg")

# require(gridExtra)
grid.arrange(mp, up, gp, tp, rp, ncol=2)
# save(rmses, file = "rmses.rda")
save(rmses, file = "rdas/rmses9.rda")
model_9_rmse <- min(rmses)
save(model_9_rmse, file = "rdas/model_9_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "9",
                                 Method="Regularized Movie + User + Genre + Rating Time + Release date"))

```

```

                    RMSE = model_9_rmse))
rmse_results <- rmse_results %>% arrange(desc(RMSE))
save(rmse_results, file = "rdas/rmse_results.rda")

# rm(train_set,test_set,movie_avgs,user_avgs,genres_avgs,time_effect_avgs,rel_effect_avgs)
# rm(lambda,lambdas, model_9_rmse,mu,rmses)
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-9")
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
Sys.time()
# load("rdas/train_set.rda")

input <- train_set

# k <- 4 # temporary, to be deleted 15-30 min
# k <- 5 # temporary, to be deleted 1 hr, also default for fold()
k <- 10 # temporary, to be deleted 3 hrs
# k <- 20 # temporary, to be deleted 5.45-6 hrs
# k <- 60 #                                19hours
Sys.time()
set.seed(1, sample.kind="Rounding") # For reproducibility

training_set <- fold(input, k = k, num_col = "rating")
# training_set <- fold(input, k = k, num_col = "rating", num_fold_cols = 2)
# training_set <- fold(input, k = k, cat_col = "genres", id_col = "userId", num_col = "rating")
# rm(input, train_set)
rm(input)

# Order by .folds
training_set <- training_set %>% arrange(.folds)

# training_set %>% kable()
str(training_set)

# lambdas <- seq(0, 10, 0.25)
# we have learnt that the lambda range is 4-7 at this point
lambdas <- seq(4, 7, 0.05)

rmses <- matrix(ncol = length(1:k), nrow = length(lambdas))
rmses <- sapply(1:k, function(fold){

  # Create training set for this iteration
  # Subset all the datapoints where .folds does not match the current fold
  train_set_ts <- training_set[training_set$.folds != fold,]

  # Create test set for this iteration
  # Subset all the datapoints where .folds matches the current fold
  test_set_ts <- training_set[training_set$.folds == fold,]
  test_set_ts <- test_set_ts %>%
    semi_join(train_set_ts, by = "movieId") %>%
    semi_join(train_set_ts, by = "userId") %>%
    semi_join(train_set_ts, by = "date")

  rmses[,fold] <- sapply(lambdas, function(l){
    mu <- mean(train_set_ts$rating)

```

```

b_i <- train_set_ts %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+1))
b_u <- train_set_ts %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+1))
b_g <- train_set_ts %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1))
b_d <- train_set_ts %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  group_by(date) %>%
  summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1))
b_r <- train_set_ts %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1))

predicted_ratings <- test_set_ts %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  left_join(b_r, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .$pred

  return(RMSE(predicted_ratings, test_set_ts$rating))
}
}
Sys.time()

rm(training_set,train_set_ts,test_set_ts)
save(rmses,file = "rdas/rmses_k.rda")

rmses_all <- as.data.frame(rmses) %>%
  mutate(lambdas=lambdas) %>%
  gather(key = "k", value = "value",1:sum(unique(k)))
rmses_all %>%
  ggplot(aes(lambdas,value, col=k)) +
  geom_point()

# target rmse 0.86490
summary(rmses)
summary(rmses_all)

```

```

# Get lambda for min rmse values for each fold
lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]
rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]

# Get mean lambda from min rmse values for each fold
mean(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])
mean(rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])
(lambda_opt <- mean(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))]))

# Get median lambda from min rmse values for each fold
median(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])
median(rmses_all$lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

# Get range/(min,max) lambda from min rmse values for each fold
range(lambdas[sapply(1:k, function(x) which.min(rmses[,x]))])

save(rmses_all,file = "rdas/rmses_all_k.rda")

Sys.time()
# load("rdas/train_set.rda")
# load("rdas/test_set.rda")

rmses_min <- sapply(lambda_opt, function(l){
  mu <- mean(train_set$rating)
  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
  b_g <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
  b_d <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    group_by(date) %>%
    summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
  b_r <- train_set %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    left_join(b_d, by = "date") %>%
    group_by(movieId) %>%
    summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

  predicted_ratings <- test_set %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_g, by = "genres") %>%
    
```

```

    left_join(b_d, by = "date") %>%
    left_join(b_r, by='movieId') %>%
    mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
    .\$pred

    return(RMSE(predicted_ratings, test_set$rating))
})

(rmses_min)
save(rmses_min, file = "rdas/rmses_min.rda")
mu <- mean(train_set$rating)
lambda <- lambda_opt
movie_reg_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n() + lambda), n_i = n())
user_reg_avgs <- train_set %>%
  left_join(movie_avgs, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n() + lambda), n_i = n())
genres_reg_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  group_by(genres) %>%
  summarize(b_g = sum(rating - mu - b_i - b_u)/(n() + lambda), n_i = n())
time_effect_reg_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  group_by(date) %>%
  summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n() + lambda), n_i = n())
rel_effect_reg_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  left_join(genres_avgs, by = "genres") %>%
  left_join(time_effect_avgs, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n() + lambda), n_i = n())

mp <- tibble(original = movie_avgs$b_i,
             regularized = movie_reg_avgs$b_i,
             n = movie_reg_avgs$n_i) %>%
  ggplot(aes(original, regularized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Movie Effect Reg")

up <- tibble(original = user_avgs$b_u,
             regularized = user_reg_avgs$b_u,
             n = user_reg_avgs$n_i) %>%
  ggplot(aes(original, regularized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("User Effect Reg")

gp <- tibble(original = genres_avgs$b_g,
             regularized = genres_reg_avgs$b_g,
             n = genres_reg_avgs$n_i) %>%
  ggplot(aes(original, regularized, size=sqrt(n))) +

```

```

geom_point(shape=1, alpha=0.5) +
  ggtitle("Genres Effect Reg")
tp <- tibble(original = time_effect_avgs$b_d,
             regularlized = time_effect_reg_avgs$b_d,
             n = time_effect_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Rating time Effect Reg")
rp <- tibble(original = rel_effect_avgs$b_r,
             regularlized = rel_effect_reg_avgs$b_r,
             n = rel_effect_reg_avgs$n_i) %>%
  ggplot(aes(original, regularlized, size=sqrt(n))) +
  geom_point(shape=1, alpha=0.5) +
  ggtitle("Release date Effect Reg")

# require(gridExtra)
grid.arrange(mp, up, gp, tp, rp, ncol=2)
model_10_rmse <- rmses_min
save(model_10_rmse, file = "rdas/model_10_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "10",
                                  Method=paste0("Regularized Movie + User + Genre + Rating Time + Release Date"),
                                  RMSE = model_10_rmse))
rmse_results <- rmse_results %>% arrange(desc(RMSE))
save(rmse_results, file = "rdas/rmse_results.rda")

# rm(train_set, test_set, movie_avgs, user_avgs, genres_avgs, time_effect_avgs, rel_effect_avgs)
# rm(lambda, lambdas, model_9_rmse, mu, rmses)
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-10")
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
load("rdas/edx_md.rda")
load("rdas/validation_md.rda")
(lambda_opt)

rmses_val <- sapply(lambda_opt, function(l){
  mu <- mean(edx_md$rating)
  b_i <- edx_md %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1), n_i = n())
  b_u <- edx_md %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1), n_i = n())
  b_g <- edx_md %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    group_by(genres) %>%
    summarize(b_g = sum(rating - mu - b_i - b_u)/(n()+1), n_i = n())
  b_d <- edx_md %>%
    left_join(b_i, by='movieId') %>%
    left_join(b_u, by='userId') %>%
    left_join(b_g, by = "genres") %>%
    group_by(date) %>%

```

```

    summarize(b_d = sum(avg_rating - mu - b_i - b_u - b_g)/(n()+1), n_i = n())
b_r <- edx_md %>%
  left_join(b_i, by='movieId') %>%
  left_join(b_u, by='userId') %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  group_by(movieId) %>%
  summarize(b_r = sum(avg_rating_rel - mu - b_i - b_u - b_g - b_d)/(n()+1), n_i = n())

predicted_ratings <- validation_md %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_g, by = "genres") %>%
  left_join(b_d, by = "date") %>%
  left_join(b_r, by='movieId') %>%
  mutate(pred = mu + b_i + b_u + b_g + b_d + b_r) %>%
  .$pred

return(RMSE(predicted_ratings, validation_md$rating))
})

(rmses_val)
save(rmses_val, file = "rdas/rmses_val.rda")
model_val_rmse <- rmses_val
save(model_val_rmse, file = "rdas/model_val_rmse.rda")

rmse_results <- bind_rows(rmse_results,
                           tibble(Index = "VAL",
                                  Method=paste0("Regularized Movie + User + Genre + Rating Time + Release",
                                                RMSE = model_val_rmse)))
rmse_results <- rmse_results %>% arrange(desc(RMSE))
save(rmse_results, file = "rdas/rmse_final_results.rda")

# rm(train_set,test_set,movie_avgs,user_avgs,genres_avgs,time_effect_avgs,rel_effect_avgs)
# rm(lambda,lambdas, model_9_rmse,mu,rmses)
kable(rmse_results, "latex", escape=FALSE, booktabs=TRUE, linesep="", caption="RMSE Results Models 1-")
  kable_styling(latex_options=c("HOLD_position"), font_size=7)
  knitr::knit_exit()
options(tinytex.verbose = TRUE)
fit <- lm(mpg ~ cyl + disp, mtcars)
# show the theoretical model
equatiomatic::extract_eq(fit)
options(tinytex.verbose = FALSE)

```

Terms like generate and some will also show up.

Alphabetical Index

```
knitr::knit_exit()
```